

BUSINESS APPLICATIONS OF AI

Module 1 — Introduction to AI Disruption

Palm Beach Atlantic University · Graduate Program

Introduction to AI Disruption

"AI doesn't disrupt one block of the Business Model Canvas. It disrupts all nine, simultaneously."

Learning Objectives

By the end of this chapter, you will be able to:

- D
isti
ng
uis
h t
he
AI
era
fro
m
pri
or t
ec
hn
olo
gic
al
wa
ve
s a
nd
ex
pla
in
wh
y L
LM
s r
epr
es
ent
a
str
uct
ura
l in
fle
cti
on,

no
t a
cyc
lica
l tr
en
d.

- A
ppl
y t
he
Bu
sin
es
s
Mo
del
C
an
va
s (
BM
C)
as
a d
isr
upt
ion
di
ag
no
stic
to
ol,
ide
ntif
yin
g
wh
ere
AI
cre
ate
s v
alu
e
or
exi
ste
nti
al r
isk
in

each of the nine blocks.

- Evaluate the strategic trade-offs of early adoption versus late adoption using the winner's curse framework.

- Configure a profile

es
sio
nal
AI
too
l st
ac
k a
nd
na
vig
ate

six
cor
e p
latf
or
ms
flu
ent
ly.

- A
nal
yz
e
a r
eal
co
mp
an
y (Sh
opi
fy)
thr
ou
gh
the
le
ns
of
AI-
nat
ive
bu
sin
es
s
mo
del
tra
nsf

or
ma
tio
n.

1.1 The Inflection

Three Eras of Business Technology

Every few decades, a technology arrives that does not merely improve how businesses operate — it rewrites what a business is. We have lived through two such moments in recent memory, and we are living through the third right now.

Era 1: Pre-Internet (1960–1994) The foundational era of business computing was defined by automation of known processes. ERP systems (SAP, Oracle) automated accounting. Relational databases (IBM DB2, Oracle) structured inventory. EDI (Electronic Data Interchange) linked suppliers to retailers. The logic was simple: take a human task, replicate it in silicon, do it faster and cheaper. The adoption timeline stretched over 20–30 years. A Fortune 500 company that bought an IBM mainframe in 1970 faced minimal urgency — competitors were on the same adoption curve. The technology was expensive, required deep IT expertise, and served as a barrier to entry rather than a leveler.

Era 2: The Internet Era (1995–2021) The internet collapsed geography and search costs. A startup in a garage could reach the same customer as a Fortune 500 company — if it could be found. Google made finding things cheap. Amazon made logistics cheap. Stripe made payments cheap. The defining dynamic was disintermediation: every layer of business that existed to mediate information (travel agents, stock brokers, classified ad publishers, encyclopedia salespeople) was compressed or eliminated. Adoption timelines shortened: a technology that took 20 years to penetrate in the pre-internet era now took 5–8 years. The winners were those who understood that software ate the world (Andreessen, 2011) before their industries did.

Era 3: The AI Era (2022–present) Something different is happening now. The release of ChatGPT in November 2022 marked not just the launch of a product, but the public legibility of a capability that had been maturing in research labs for a decade. What changed was not compute power alone — it was interpretability. For the first time, a machine could receive an ambiguous, natural-language instruction and return a useful, contextually appropriate response. This is not retrieval. It is not pattern-matching in the narrow sense. It is interpretation.

Why LLMs Are Different: Interpretation, Not Retrieval

To understand why LLMs represent a structural break from prior AI systems, consider the following comparison:

⌋ Prior AI Systems (1980–2021) Examples: Spam filters, recommendation engines, fraud detection, image classifiers, search ranking

How they work:

- Tr

ained
ed
on
lab
ele
d d
ata
set
s w
ith
cle
ar r
igh
t/w
ron
g a
ns
we
rs

- Operate within a closed problem space defined at training time

- Retrieve or classify by based

on
sta
tisti
cal
si
mil
arit
y

- Fail
grace
fully
: un
cer
tain
=
no
action
or
es
cal
ate
to
hu
ma
n

- Re
quire
de
fine
d
-sp
eci
fic
retr
ain
ing

for
ea
ch
ne
w t
as
k

Business implication: Automate repetitive decisions within known problem spaces. Replace or augment specific workers doing specific tasks. High cost per deployment, high specificity.

This distinction matters enormously for business strategy. When AI could only automate known, structured tasks, a company's competitive moat was still built on proprietary knowledge, relationships, and brand. When AI can reason — when it can synthesize a market analysis, draft a contract, write production code, plan a supply chain restructuring, and explain all of it in plain language — the moat shifts. Proprietary data becomes the moat, not the knowledge workers who interpret it.

The Compression of Adoption Timelines

History shows a consistent pattern: each technological wave compresses adoption timelines relative to the prior wave. The printing press took roughly 50 years to achieve widespread European adoption. Electricity took 40 years to penetrate American manufacturing. The telephone took 35 years to reach 50% of US households. The internet took approximately 7 years to reach 50% of the US population after the commercialization of the web browser. Smartphones reached 50% penetration in about 5 years.

ChatGPT reached 100 million users in two months — the fastest consumer technology adoption in history. By 2025, Andreessen Horowitz's State of AI report found that AI tool usage had penetrated 40% of knowledge workers globally, with enterprise AI budgets growing at 120% year-over-year.

What does this compression mean for business strategy? Simply: the window for strategic positioning is shorter than it has ever been. In the internet era, a company that recognized the threat in 2000 had until approximately 2005–2007 to build a response before incumbents were structurally disadvantaged. In the AI era, the equivalent window may be 12–24 months.

This is not hyperbole. Consider:

- KI
arn
a d
epl
oy
ed
AI
to r
epl
ac
e 7
00
cu
sto
me
r s
erv
ice
ag
ent
s
in
20
24.
Wi
thi

n
12
mo
nth
s,
it r
ep
ort
ed
tha
t
its
AI
sys
te
m
ha
ndl
ed
2/3
of
cu
sto
me
r s
erv
ice
ch
ats
wit
h s
ati
sfa
cti
on
sc
ore
s
ma
tch
ing
hu
ma
n a
ge
nts

—
at
1/1
0th
th
e c

ost

.

- G
old
ma
n S
ac
hs
rep
ort
ed
in
20
25
tha
t
AI
co
din
g a
ssi
sta
nts
ha
d r
ed
uc
ed
sof
tw
are
de
vel
op
me
nt t
im
e
by
40
%
for
its
en
gin
eer
ing
te
am
s,
eff
ect
ivel

y e
xp
an
din
g
its
en
gin
eer
ing
ou
tpu
t wi
tho
ut
he
ad
co
unt
gr
ow
th.

- M
cKi
ns
ey'
s 2
02
5 S
tat
e
of
AI
rep
ort
fou
nd
tha
t c
om
pa
nie
s
in t
he
top
qu
arti
le
of
AI
ad
opt

ion
w
ere
ac
hie
vin
g r
ev
en
ue
gro
wt
h 1
.4x
hi
gh
er t
ha
n t
hei
r b
ott
om
-qu
arti
le
pe
ers

—
an
d t
hat
ga
p
wa
s w
ide
nin
g.

The Winner's Curse

The compression of adoption timelines creates a strategic paradox known in this context as the winner's curse of AI adoption. The term borrows from auction theory, where the winner of an auction tends to have overpaid because they have the highest estimate of value — often because they are also the most optimistic.

In the AI context, the winner's curse operates differently: the companies that move first face the highest implementation costs and the lowest model quality, while the companies that move last face the highest competitive disadvantage. Both extremes destroy value.

The optimal adoption strategy is neither first-mover nor fast-follower in the classical sense. It is

informed early majority: move after the technology has achieved sufficient reliability for production use, but before the late majority normalizes it. In 2026, we are at or approaching this inflection for most enterprise AI applications.

1.2 The Business Model Canvas as a Disruption Map

Osterwalder and Pigneur's Business Model Canvas (2010) is the dominant framework for visualizing and analyzing business models. Originally designed as a tool for design and iteration, it has become, in the AI era, an indispensable diagnostic for disruption. Each of the nine blocks represents a surface where AI is simultaneously creating new value and destroying old competitive positions.

The canvas framework is deceptively simple. Behind each block lies a complex set of assumptions about how value is created, delivered, and captured. AI disrupts these assumptions — sometimes subtly, sometimes catastrophically.

```
graph TD
    A[Business Model Canvas] --> B[Customer Segments]
    A --> C[Value Propositions]
    A --> D[Channels]
    A --> E[Customer Relationships]
    A --> F[Revenue Streams]
    A --> G[Key Resources]
    A --> H[Key Activities]
    A --> I[Key Partnerships]
    A --> J[Cost Structure]

    B --> B1["AI Impact: Segment-of-one\nhyper-personalization"]
    C --> C1["AI Impact: Dynamic value props\nco-created with AI"]
    D --> D1["AI Impact: Conversational\nchannels replace browse-first"]
    E --> E1["AI Impact: Always-on\npersonalized agents"]
    F --> F1["AI Impact: Usage-based\noutcome-based pricing"]
    G --> G1["AI Impact: Data and models\nbecome primary resources"]
    H --> H1["AI Impact: AI-augmented\nknowledge work at scale"]
    I --> I1["AI Impact: AI vendors\nbecome critical partners"]
    J --> J1["AI Impact: Marginal cost\nof knowledge work !' 0"]
```

Block 1: Customer Segments

How AI disrupts it. AI is collapsing the granularity of segmentation toward its theoretical limit: the segment of one. Traditional segmentation grouped customers because personalization was expensive — it required human labor to customize offers, content, and communication for each individual. AI makes this labor negligible.

Current examples (2025–2026):

- Netflix's AI per

so
nal
iza
tio
n e
ngi
ne
no
w
ge
ner
ate
s i
ndi
vid
ual
ize
d t
hu
mb
nai
l ar
tw
ork
, tr
ail
er
cut
s,
an
d c
ont
ent
re
co
m
me
nd
ati
on
s t
hat
va
ry
by
us
er,
de
vic
e, t
im
e
of

da
y, a
nd
em
oti
on
al
sta
te i
nfe
rre
d fr
om
vi
ewi
ng
pat
ter
ns.
Th
e r
es
ult:
a
me
as
ura
ble
re
du
cti
on
in
ch
urn
an
d
a 3
5%
in
cre
as
e
in
co
nte
nt
dis
co
ver
y p
er
se

ssi
on.

- S
pot
ify'
s
DJ
fea
tur
e u
se
s
AI
to
ge
ner
ate
pe
rso
nal
ize
d a
udi
o c
om
me
nta
ry
bet
we
en
so
ng
s,
cre
ati
ng
the
ex
per
ien
ce
of
a r
adi
o
DJ
wh
o k
no
ws
the
list

en
er i
nti
ma
tel
y. T
his
fe
atu
re
dro
ve
a 1
2%
in
cre
as
e
in
pre
mi
um
su
bs
cri
pti
on
s a
mo
ng
us
ers
w
ho
en
ga
ge
d
wit
h
it
in
20
24.
• S
titc
h F
ix
ha
s u
se
d
Al-

drive
n s
eg
me
nt-
of-
on
e s
tyli
ng
sin
ce
20
19,
bu
t
its
20
25
"St
yle
AI"
sy
ste
m
go
es
furt
her
: sy
nth
esi
zin
g p
urc
ha
se
his
tor
y, s
oci
al
me
dia
pr
efe
ren
ce
s i
nfe
rre
d fr

om
br
ow
sin
g d
ata
, a
nd
se
as
on
al
co
nte
xt
to
pro
du
ce
styl
e r
ec
om
me
nd
ati
on
s t
hat
be
at
hu
ma
n s
tyli
sts

on
cu
sto
me
r s
ati
sfa
cti
on
sc
ore
s.

Strategic implication: Companies that continue to operate with coarse segmentation (e.g., "millennial women, income \$50K–\$80K") are leaving conversion and retention on the table. The segmentation advantage now belongs to those with the richest first-party data and the AI infrastructure to activate

it.

Block 2: Value Propositions

Traditional definition: The bundle of products and services that create value for a specific Customer Segment. The value proposition answers the question: why should a customer choose us?

How AI disrupts it: AI creates value propositions that were previously impossible: propositions that are dynamic (changing in real-time based on context), co-created (built in dialogue with the customer), and compound (aggregating capabilities that previously required multiple vendors).

Current examples:

• Harvard
AI (legal)
of
fer
s
a v
alu
e p
rop
osi
tio
n i
mp
os
sibl
e b
efo
re
LL
Ms
: a
jun
ior
as
so
cia
te-l
ev
el l
eg
al r
es
ear
ch
an
d d

raft
ing
ca
pa
bilit
y a
vail
abl
e 2
4/7
at
\$0.
01
per
ta
sk.
Thi
s
is
not
fa
ste
r le
gal
re
se
arc
h
—
it
is
a n
ew
va
lue
pr
op
osi
tio
n
for
leg
al
wo
rk t
hat
di
spl
ac
es
the
ho
url

y b
illin
g
mo
del

.
• G
itH
ub
Co
pil
ot
off
ers
"p
air
pro
gra
m
mi
ng
wit
h
an
AI t
hat
kn
ow
s e
ver
y p
ubl
ic c
od
eb
as
e."
In
20
25,
Gi
tH
ub
rep
ort
ed
tha
t C
opi
lot
us
ers
co

mp
let
ed
tas
ks
55
% f
ast
er
an
d
wr
ote
co
de
wit
h 4
0%
fe
we
r b
ug
s.
Thi
s
is
not

an
inc
re
me
nta
l i
mp
rov
em
ent

—
it
is
a s
tep
-ch
an
ge
in
wh
at
de
vel
op

er
pro
du
ctiv
ity
me
an
s.

- G
lea
n (ent
erp
ris
e s
ear
ch)
pr
om
ise
s t
hat
ev
ery
e
mp
loy
ee
ca
n fi
nd
an
y i
nte
rna
l in
for
ma
tio
n
in
se
co
nd
s.
Thi
s p
rop
osi
tio
n
wa
s t

ec
hni
call
y a
chi
ev
abl
e b
efo
re
LL
Ms
bu
t e
co
no
mi
call
y u
nvi
abl
e.
Wit
h
AI,
Gl
ea
n's
in
de
xin
g a
nd
retr
iev
al i
nte
llig
en
ce
ma
ke
s t
he
pro
po
siti
on
bot
h a
chi
ev
abl

e a
nd
aff
ord
abl
e.

Strategic implication: Audit your current value proposition against this question: Is this proposition becoming more or less differentiated as AI capabilities improve? If a competitor can now deliver your core value proposition with AI at 1/10th the cost, your moat is eroding.

Block 3: Channels

Traditional definition: The touchpoints through which a company communicates with and reaches its Customer Segments to deliver a Value Proposition.

How AI disrupts it: AI is fundamentally changing the nature of channels, not just their efficiency. The shift is from browse-first channels (websites, stores, app catalogs) to conversational-first channels (AI agents, voice interfaces, chat-native experiences).

Current examples:

- P
erp
lexi
ty
AI
ha
s g
ro
wn
to
20
mill
ion
da
ily
act
ive
us
ers

by
20
25
as
an
AI-
nat
ive
se
arc
h c

ha
nn
el.
Br
an
ds
tha
t o
pti
mi
ze
d
for
Go
ogl
e S
EO
no
w f
ac
e
a n
ew
ch
an
nel
:
Al
an
sw
er
en
gin
es
tha
t s
ynt
he
siz
e i
nfo
rm
ati
on
rat
her
th
an
list
link
s.
Wi
nni

ng
in t
his
ch
an
nel
re
qui
res
be
ing
cit
ed
by
AI,
not
ju
st r
an
ke
d
by
alg
orit
hm
s.

- S
ho
pif
y's
AI
sto
refr
ont
as
sist
ant
all
ow
s c
ust
om
ers
to
de
scr
ibe
w
hat
th
ey
wa
nt

in
nat
ura
l la
ng
ua
ge
("I
ne
ed
a b
irth
da
y g
ift
for
my
sis
ter
wh
o l
ov
es
hiki
ng,
bu
dg
et
\$5
0—
\$8
0")
an
d r
ec
eiv
e c
ura
ted
re
co
m
me
nd
ati
on
s.
Thi
s c
oll
ap
se
s t

he
bro
ws
e-a
nd-
filt
er
UX
int
o
a c
on
ver
sat
ion
al
ex
per
ien
ce.

- D
uol
ing
o's
Al
tut
or (
Ma
x) r
epl
ac
ed
mu
ch
of t
he
ap
p-b
as
ed
les
so
n s
tru
ctu
re
wit
h c
on
ver
sat
ion
al

practic
e.
Engage
ment
impe
rse
ssi
on
increa
sed 3
0%

as
users
found
the
con
versations
natural
and
effective.

Strategic implication: Your channel strategy must now include AI-mediated channels. This means optimizing for AI discoverability (structured data, authoritative content), building conversational interfaces, and understanding how AI agents make recommendations on behalf of users.

Block 4: Customer Relationships

Traditional definition: The types of relationships a company establishes with specific Customer Segments — from self-service to dedicated personal assistance, from transactional to community-based.

How AI disrupts it: AI enables a new category of customer relationship: the always-on, personalized AI companion that combines the warmth of a dedicated personal advisor with the scale of self-service. This collapses the traditional trade-off between relationship depth and operational cost.

Current examples:

- Klarna's AI assistant serves as a shopping advisor, financial manager, and customer service agent simultaneously.

In 2024, it handled 2.3 million customer conversations in a single day — equivalent to the workload of 5,000 human agents. Customer satisfaction

score
were
within
2
points
of
human
agent
performance.

- Scale
force
Einst
ein
AI
provid
es
sales
es
represent
ative
switch
ith
real
l-time
rel
ation
ship
intelligence:
fla

ggi
ng
at-r
isk
ac
co
unt
s, r
ec
om
me
ndi
ng
ne
xt
be
st
act
ion
s,
an
d d
raft
ing
pe
rso
nal
ize
d o
utr
ea
ch.
Th
e r
es
ult
is
not
ju
st
effi
cie
nc
y
—
it
is
a q
ual
itat
ivel
y d
iffe

relationship between the sales representative, the AI, and the customer.

- Webot (mental health) has demonstrated that AI companion can maintain therapy

uti
c-q
ual
ity
rel
ati
on
shi
ps
wit
h u
ser
s e
xp
eri
en
cin
g
mil
d a
nxi
ety
an
d d
epr
es
sio
n,
wit
h c
lini
cal
ou
tco
me
s c
om
par
abl
e
to
CB
T
for
sel
ect
co
ndi
tio
ns
(Fit
zp
atri

ck
et
al.,
20
17)

.
By
20
25,
W
oe
bot
se
rve
d o
ver
4
mill
ion
us
ers

—
a s
cal
e
no
hu
ma
n t
her
api
st
wo
rkf
orc
e c
oul
d
ma
tch
.

Block 5: Revenue Streams

Traditional definition: The cash a company generates from each Customer Segment. Revenue models include asset sales, usage fees, subscription fees, licensing, brokerage fees, and advertising.

How AI disrupts it: AI is enabling entirely new revenue architectures built around outcomes rather than inputs, and usage rather than access. The shift is from "pay for the software" to "pay for the result."

Current examples:

• E
ve
nU
p (l
eg
al
AI)
ch
arg
es
law
fir
ms
a
per
ce
nta
ge
of
set
tle
me
nts
se
cur
ed
usi
ng
its
AI-
ge
ner
ate
d d
em
an
d l
ett
ers
. T
his
ou
tco
me
-ba
se
d
mo
del
w
as

only possible because even 'Up' is AI measureably predictive of settlement outcomes. They price on value, not inputs.

- Otter.ai shifted from a flat subs

cri
pti
on
to
a u
sa
ge-
ba
se
d
mo
del
in
20
24,
ch
arg
ing
pe
r m
eet
ing
tra
ns
cri
be
d.
As
tra
ns
cri
pti
on
be
ca
me
a
co
m
mo
dit
y, t
he
diff
ere
nti
ati
on
mo
ve
d
to
an

analysis and action item extraction

— each billed as a premium add-on.

- Palantir's AI Platform (AIP) uses a hybrid model: a base subscription for

m
ac
ce
ss
plu
s o
utc
om
e-li
nk
ed
bo
nu
se
s
for
me
as
ura
ble
op
era
tio
nal
im
pro
ve
me
nts
. F
or
a d
efe
ns
e c
ont
rac
tor
clie
nt
in
20
25,
thi
s tr
an
sla
ted
to
\$4
0M
in
out

comes on a \$10M platform minus subscriptions.

Strategic implication: Evaluate whether your current pricing model is aligned with the value AI is actually creating for your customers. If your customers are capturing AI-generated value that you are not pricing, you are leaving money on the table.

Block 6: Key Resources

Traditional definition: The most important assets required to make a business model work — physical, intellectual, human, or financial.

How AI disrupts it: Data and AI models are emerging as the dominant Key Resources in the AI era, displacing human expertise and physical assets in many business models. The competitive moat is shifting from "we have the best people" to "we have the best data and the best models trained on it."

Current examples:

- Bloomberg GPT (2023) demonstrated that a dominant specific

fic
LL
M t
rai
ne
d
on
pro
pri
eta
ry f
ina
nci
al
dat
a o
utp
erf
or
me
d g
en
era
l-p
urp
os
e L
LM
s
on
fin
an
cial
ta
sks
. Bl
oo
mb
erg
's
de
ca
de
s
of
str
uct
ure
d fi
na
nci
al
dat

a b
ec
am
e
a c
om
pet
itiv
e r
es
our
ce
for
AI
de
vel
op
me
nt.

• V
ee
va
Sy
ste
ms
(lif
e s
cie
nc
es
Sa
aS
) b
uilt
a
co
mp
etit
ive
m
oat
ar
ou
nd
clin
ical
tri
al
dat
a a
nd
reg
ula

tor
y d
oc
um
ent
ati
on.

Its
AI l
ay
er,
trai
ne
d
on
thi
s p
rop
riet
ary
co
rpu
s,
no
w
pro
vid
es
pre
dic
tiv
e a
nal
ytic
s
on
dru
g a
ppr
ov
al
pro
ba
bilit
y
—
a c
ap
abi
lity
no
co

mp
etit
or
ca
n r
epl
ica
te
wit
ho
ut t
he
dat
a.

- D
ee
pM
ind
's
Alp
ha
Fol
d d
em
on
str
ate
d t
hat
Al
mo
del
s,
on
ce
trai
ne
d,
be
co
me
int
ell
ect
ual
pr
op
ert
y
of
ext
rao
rdi

nar
y v
alu
e.
Ast
raZ
en
ec
a li
ce
ns
ed
Alp
ha
Fol
d a
cc
es
s
to
ac
cel
era
te
dru
g d
isc
ov
ery

—
pa
yin
g
for
the
m
od
el,
not
th
e p
eo
ple
w
ho
bui
lt
it.

Strategic implication: Audit your Key Resources with fresh eyes: what data do you generate that is proprietary, high-quality, and large-scale? That data is likely your most important AI resource. If you are not capturing, labeling, and securing it, you are not building an AI-era competitive moat.

Block 7: Key Activities

Traditional definition: The most important things a company must do to make its business model work — production, problem solving, platform/network management.

How AI disrupts it: AI is restructuring what "doing" means. Activities that once required large teams of skilled workers are now AI-augmented or AI-executed. The key activity is increasingly the design and management of AI systems, rather than the execution of the underlying tasks those systems perform.

Current examples:

- J
as
per
AI
(co
nte
nt
ma
rke
tin
g)
shi
fte
d
its
key
ac
tivit
y fr
om
co
nte
nt
cre
ati
on
to
co
nte
nt
str
ate
gy
an
d
AI
orc
he
str
ati
on.

Its
pla
tfor
m
all
ow
s
ma
rke
tin
g t
ea
ms
to
def
ine
br
an
d v
oic
e,
ge
ner
ate
co
nte
nt
at
sc
ale
, a
nd
opt
imi
ze
for
per
for
ma
nc
e
—
wit
h
mi
ni
ma
l h
um
an
wri
tin

g i
nv
olv
ed
in
ex
ec
uti
on.

• L
eg
alz
oo
m'
s
Al
do
cu
me
nt
ge
ner
ati
on
shi
fte
d t
he
key
ac
tivit
y
of l
eg
al
do
cu
me
nt
cre
ati
on
fro
m l
ice
ns
ed
att
orn
ey
s
to
Al

sys
te
ms
, wi
th
att
orn
ey
s f
oc
usi
ng
on
ed
ge
ca
se
s a
nd
qu
alit
y a
ss
ura
nc
e.
Thi
s r
est
ruc
tur
ed
bot
h c
ost
str
uct
ure
an
d h
um
an
ca
pit
al r
eq
uir
em
ent
s.

- Wa
ym

o (autonomous vehicles) illustrates the extreme end: the entire key activity of driving has been delegated to AI systems, with human key activities reduced

to f
lee
t m
an
ag
em
ent
, s
afe
ty
mo
nit
ori
ng,
an
d r
eg
ula
tor
y c
om
pli
an
ce.

Block 8: Key Partnerships

Traditional definition: The network of suppliers and partners that make the business model work. Partnerships range from strategic alliances to supplier relationships to joint ventures.

How AI disrupts it: A new category of critical partner has emerged: AI foundation model providers. OpenAI, Anthropic, Google, and Meta have become infrastructure-layer partners for thousands of businesses — with all the strategic risks that infrastructure dependency implies.

Current examples:

- C
urs
or (
AI
co
de
edi
tor)
bu
ilt
its
ent
ire
pro
du
ct
on

top
of
Ant
hro
pic'
s C
lau
de
AP
I.
Wh
en
Ant
hro
pic
ch
an
ge
d
its
AP
I pr
icin
g
in
20
24,
C
urs
or'
s u
nit
ec
on
om
ics
we
re
dir
ect
ly i
mp
act
ed
— i
llus
trat
ing
th
e p
latf
or
m r

risk of a single AI partner dependency.

- In the context of integrated multiple AI providers (OpenAI, Anthropic) into its customer service platform, explicitly to reduce

uc
e s
ing
le-
pro
vid
er r
isk.

By
20
25,
thi
s
mu
lti-
mo
del
ar
chi
tec
tur
e h
ad
be
co
me
a
sta
nd
ard
re
co
m
me
nd
ati
on
for
ent
erp
ris
e
AI
pro
cur
em
ent
.

- M
icr
os
oft'

s
Op
en
AI
par
tne
rsh
ip (
\$1
3B
inv
est
me
nt)
is t
he
def
ini
ng
ex
am
ple
of
AI
par
tne
rsh
ips
re
sh
api
ng
co
mp
etit
ive
po
siti
on
s.
Mi
cro
sof
t's i
nte
gra
tio
n
of
GP
T-4
ac
ros

s
Off
ice
36
5,
Git
Hu
b,
Az
ure
, a
nd
Bin
g c
rea
ted
a
co
mp
ou
ndi
ng
co
mp
etit
ive
ad
va
nta
ge
tha
t G
oo
gle
, A
ma
zo
n,
an
d S
ale
sfo
rce
sc
ra
mb
led
to
rep
lica
te
wit
h t

heir own AI partner ecosystem.

Block 9: Cost Structure

Traditional definition: The costs incurred to operate a business model — fixed costs, variable costs, economies of scale, economies of scope.

How AI disrupts it: AI is driving the marginal cost of knowledge work toward zero. This is the most structurally disruptive force in the entire canvas, because it collapses the cost advantage of scale that large incumbents have historically enjoyed.

Current examples:

- Doubling or reduced its contract or translation workload by 10% in 2024, explicitly

ng
AI
as
the
re
pla
ce
me
nt.
Th
e c
ost
of
co
nte
nt l
oc
aliz
ati
on
—
a h
ist
ori
call
y e
xp
en
siv
e,
hu
ma
n-i
nte
nsi
ve
act
ivit
y
—
coll
ap
se
d.
Th
e s
avi
ng
s w
ere
rei
nv
est

ed
in
AI
pro
du
ct
de
vel
op
me
nt.

• I
BM
an
no
un
ce
d
a p
au
se
on
hiri
ng
for
7,8
00
po
siti
on
s e
xp
ect
ed
to
be
rep
lac
ed
by
AI
in
20
23.

By
20
25,
IB
Mr
ep
ort
ed

that AI had replaced 40% of its HR functions with a 25% improvement in processing speed and a 30% reduction in error rates.

- The gg reported a 48%

re
ve
nu
e d
ecli
ne
in
20
23
–2
02
4,
dir
ect
ly
attr
ibu
tin
g t
he
los
s
to
stu
de
nts
us
ing
C
hat
GP
T i
nst
ea
d
of
its
tut
ori
ng
su
bs
cri
pti
on.
Th
e c
ost
str
uct
ure
of
hu

making it
possible to
compete
with the
nearest
rival
on a
marginal
cost of
AI that
is
g.

Strategic implication: The zero-marginal-cost economics of AI knowledge work mean that cost leadership strategies based on labor arbitrage are increasingly fragile. If your cost advantage is built on cheap knowledge workers (offshore outsourcing, large analyst teams), model the scenario where AI reduces that cost to near zero for your competitors.

1.3 The Tool Stack / Lab Setup

This course is built around a philosophy: you cannot teach AI business strategy from the outside. You must use the tools, internalize their capabilities and limitations, and develop intuitions that can only come from hands-on practice.

The following six tools constitute your primary workbench. All are free or have generous free tiers. All are used by professional AI practitioners in 2025–2026.

```

flowchart LR
    subgraph Core_Reasoning ["Core Reasoning"]
        A["gemini.google.com\nConversational AI\nMultimodal reasoning"]
        B["claude.ai\nAdvanced reasoning\nLong-context analysis"]
    end
    subgraph Knowledge ["Knowledge & Research"]
        C["notebooklm.google.com\nDocument synthesis\nSource-grounded answers"]
    end
    subgraph Development ["Development & Experimentation"]
        D["aistudio.google.com\nAPI access & prompting\nModel comparison"]
        E["Claude Code\nAI-native coding\nFull codebase awareness"]
    end
    subgraph Collaboration ["Creative & Collaboration"]
        F["Google Antigravity\nTeam AI workspace\nShared context"]
    end

    A --> C
    B --> E
    D --> E
    C --> F

```

Tool 1: gemini.google.com

uni
 mo
 dal
 re
 as
 oni
 ng:
 G
 em
 ini
 ca
 n a
 nal
 yz
 e i
 ma
 ge
 s,
 PD
 Fs,
 sp
 rea
 ds
 he
 ets
 , a
 nd
 vid
 eo
 alo

ng
sid
e t
ext
. U
plo
ad
a c
om
pet
itor
's
an
nu
al r
ep
ort
an
d a
sk
Ge
mi
ni
to i
de
ntif
y t
hei
r
AI i
nv
est
me
nt
sig
nal
s.

- D
ee
p
Re
se
arc
h
mo
de:
G
em
ini'
s D
ee
p
Re

se
arc
h f
eat
ure
co
nd
uct
s
mu
lti-
ste
p
we
b r
es
ear
ch
an
d s
ynt
he
siz
es
fin
din
gs
int
o s
tru
ctu
red
re
por
ts
—
rep
laci
ng
4—
6 h
our
s
of
an
aly
st
wo
rk
in
20
—
30

minutes.

• Google Workspace integration: Gemini is embedded in Docs, Sheets, Gmail, and Slides, allowing AI assistance within your existing workflow.

ow.

Best use cases in this course: Market research, competitive analysis, document summarization, brainstorming, drafting.

Tool 2: notebooklm.google.com

NotebookLM is Google's document-grounded AI workspace. Unlike general-purpose chatbots, NotebookLM is explicitly anchored to documents you upload. It will not hallucinate facts outside your source material — and it will cite the page and paragraph for every claim.

Best use cases in this course: Analyzing academic papers, synthesizing primary research, preparing for case study discussions, building literature reviews.

Tool 3: aistudio.google.com

Google AI Studio is the developer-facing interface for Google's Gemini family of models. It provides:

- Prompt engineering workspace: Test prompts, compare outputs, and refine systematically.
- M

od
el
sel
ect
ion
: C
om
par
e
Ge
mi
ni
Pr
o,
Fla
sh,
an
d s
pe
cial
ize
d v
ari
ant
s.

- A
PI
key
m
an
ag
em
ent
: G
en
era
te
AP
I k
ey
s
for
pro
gra
m
ma
tic
ac
ce
ss.

- G
rou

ndi
ng
wit
h
Go
ogl
e S
ear
ch:
En
abl
e r
eal
-ti
me
w
eb
se
arc
h a
ug
me
nta
tio
n
for
tim
e-s
en
siti
ve
qu
eri
es.

Best use cases in this course: Prompt engineering, API experimentation, building custom AI workflows, testing model performance on business tasks.

Tool 4: claude.ai

Anthropic's Claude is, as of 2025–2026, the most capable AI system for complex reasoning, nuanced writing, and long-context analysis. Claude's distinguishing features:

- E
xte
nd
ed
thi
nki
ng:
Cl
au
de'

se
xte
nd
ed
thi
nki
ng
mo
de
sh
ow
s
its
rea
so
nin
g p
roc
es
s,
en
abl
ing
yo
u
to
au
dit
an
d g
uid
e
its l
ogi
c.

• 2
00
K t
ok
en
co
nte
xt
win
do
w:
Cla
ud
e c
an
hol
d
an

ent
ire
bo
ok,
le
gal
co
ntr
act

,
or f
ina
nci
al
mo
del
in
co
nte
xt
si
mu
lta
ne
ou
sly.

• C
on
stit
uti
on
al
AI:
Ant
hro
pic'
s s
afe
ty t
rai
nin
g
ma
ke
s C
lau
de
not
abl
y
mo
re
car

efu
l a
bo
ut
un
cer
tai
nty,
m
ore
lik
ely
to
ex
pre
ss
do
ubt
w
he
n a
ppr
opr
iat
e,
an
d l
es
s p
ron
e
to
co
nfi
de
nt
co
nfa
bul
ati
on.

Best use cases in this course: Complex analysis, long-document work, nuanced strategic reasoning, writing with high accuracy requirements.

Tool 5: Claude Code

Claude Code is Anthropic's AI-native coding agent — distinct from `claude.ai` in that it operates as a terminal-based agent with full access to your codebase, file system, and development environment.

What this means in practice:

- C

lau
de
Co
de
ca
n r
ea
d,
wri
te,
an
d e
xe
cut
e c
od
e a
cro
ss
yo
ur
ent
ire
pro
jec
t

- It
ca
n i
nst
all
de
pe
nd
en
cie
s, r
un
tes
ts,
de
bu
g e
rro
rs,
an
d r
efa
cto
r ar
chi
tec
tur

e
• It
op
era
tes
au
ton
om
ou
sly
on
mu
lti-
ste
p t
as
ks:
"bu
ild
a
we
b s
cra
per

for
co
mp
etit
or
pri
cin
g" !'
run
s
for
10
mi
nut
es
!' ret
urn
s w
ork
ing
co
de

Best use cases in this course: Lab assignments with coding components, building MVPs, data analysis automation.

Tool 6: Google Antigravity

Google Antigravity (formerly known as Google Labs' Workspace AI collaboration tools) represents Google's vision for AI-native collaborative work. It enables:

- Shared AI context: Multiple team members working with a shared AI that has a context on the team's collective knowledge
- Workflow au

to
ma
tio
n:
Co
nn
ect
ing
AI
rea
so
nin
g
to
do
cu
me
nt
cre
ati
on,
e
ma
il, c
ale
nd
ar,
an
d p
roj
ect
m
an
ag
em
ent
• E
nte
rpri
se-
gra
de
pri
va
cy:
All
dat
a p
roc
es
se
d
wit

hin
G
oo
gle
's
ent
erp
ris
e s
ec
urit
y b
ou
nd
ary

Best use cases in this course: Team projects, capstone work, building shared knowledge bases.

Supporting Tools

Tool	Purpose	Free Tier
[GitHub](https://github.com)	Code repository, version control, deployment	Free for public repos
[Cursor](https://www.cursor.com)	AI-native code editor (VS Code with AI)	2-week trial, then \$20/mo
[Obsidian](https://obsidian.md)	Personal knowledge management, linked notes	Free for personal use

Case Study: Shopify's Transformation from Commerce Platform to AI-Native Merchant OS

Background

Shopify was founded in 2006 by Tobias Lütke, who built the platform out of frustration with existing e-commerce software while trying to sell snowboards online. By 2021, Shopify powered 1.7 million merchants globally and had a market capitalization of \$185 billion — briefly making Lütke one of the wealthiest people in Canada.

Then came 2022. Post-pandemic normalization of e-commerce growth, rising interest rates, and competition from Amazon, TikTok Shop, and Wix compressed Shopify's market cap to \$35 billion by late 2022. The company was forced to lay off 10% of its workforce and confront an existential question: What is Shopify's value proposition in a world where AI can do much of what Shopify's human services do?

Lütke's answer was not defensive. In a 2023 letter to employees, he wrote: "Shopify Sidekick is the most important product we've ever built." Sidekick was not a feature — it was a strategic repositioning.

The BMC Analysis: Before and After

Before (2015–2021): The Commerce Platform

Block	Description
Customer Segments	Small-to-medium merchants; direct-to-consumer brands
Value Proposition	"Commerce made easy" — turnkey online store with payments, inventory, shipping
Channels	SEO/content marketing, partner app marketplace, word of mouth
Customer Relationships	Self-service SaaS; Shopify Plus for enterprise (dedicated support)
Revenue Streams	Monthly subscription tiers; Shopify Payments transaction fees
Key Resources	Platform infrastructure; payment network; app ecosystem (7,000+ apps)
Key Activities	Platform development; merchant acquisition; payment processing
Key Partnerships	Payment processors; shipping carriers; app developers
Cost Structure	Infrastructure; R&D; S&M at 30–40% of revenue

After (2023–2026): The AI-Native Merchant OS

Block	Description
Customer Segments	Same base + new segments: enterprise merchants requiring AI capabilities
Value Proposition	"AI that runs your store" — Sidekick handles marketing copy, inventory forecasting, customer segmentation
Channels	Conversational AI onboarding; AI-generated storefront personalization
Customer Relationships	Sidekick as always-on AI business advisor for every merchant
Revenue Streams	New: AI add-on subscriptions; outcome-based pricing pilots for AI-driven revenue uplift
Key Resources	Merchant transaction data (billions of data points on purchasing behavior); AI models trained on this corpus
Key Activities	AI model training on proprietary merchant data; Sidekick capability expansion
Key Partnerships	OpenAI (Sidekick engine); Google (search and shopping integrations); logistics AI partners
Cost Structure	Shift: infrastructure costs rising (AI inference); headcount costs falling (AI replaces support and content roles)

The Strategic Insight

Shopify's most defensible asset was never its software — it was the behavioral data generated by 2+ million merchants and their customers. This data, accumulated over 17 years, represents purchasing patterns, conversion rates, inventory dynamics, and marketing effectiveness across virtually every product category.

By training AI models on this proprietary corpus, Shopify created a moat that a startup with equivalent AI talent cannot replicate without equivalent data. The AI capability is the data — and the data is the competitive advantage.

Key lessons from Shopify's transformation:

- In
cu
mb
ent
s w
ith
dat
a
mo
ats
ca
n
win
th
e
AI t
ran
siti
on
—
if t
he
y
mo
ve
bef
ore
th
eir
dat
a a
dv
ant
ag
e e
rod
es.

- AI
rep
osi
tio
ns
exi
sti
ng
res
our
ce
s
as

competitive advantages

— Shopify's transaction data was invaluable before AI; it is now the primary competitive resource.

- The value proposition shifts

from
multitool
to
advisor
—
Sho
opify
went
from
her
e's
a
platform
to
her
e's

an
Alt
hat
helps
your
business
grow."

- Revenue
model
evolu
tion
follow
scap
abi

lity
ev
olu
tio
n
—
as
Sid
eki
ck
de
mo
nst
rab
ly
dri
ve
s
me
rch
ant
re
ve
nu
e,
out
co
me
-ba
se
d p
rici
ng
be
co
me
s vi
abl
e.

Faith Integration — Module 1

"The Lord God took the man and put him in the Garden of Eden to work it and take care of it."
— Genesis 2:15 (NIV)

Stewardship in the Age of AI

From the opening chapters of Scripture, humans are called to be stewards — caretakers of what God has entrusted to them. We did not create the earth; we were placed in it with a mandate to cultivate and keep it. The same logic applies to every resource, talent, and technology that passes through our hands.

Artificial intelligence is one of the most powerful tools ever placed in human hands. Like fire, the printing press, or the internet before it, it does not arrive with a pre-programmed moral direction — it amplifies the intentions and values of those who wield it. A Christian business professional operating in this environment faces a question that goes beyond market strategy: Am I using this tool to serve, or to exploit? To build up, or to extract?

The stewardship framework gives us a lens for evaluating AI adoption decisions that return-on-investment calculations cannot. It asks: Does this application of AI reflect care for people, for communities, for the world God made? Does it leave things better than we found them? Does it align the power of technology with the purposes of love?

As you configure your AI tool stack in this module, resist the temptation to treat these tools as mere productivity multipliers. They are instruments of agency — your agency. The character you bring to them will shape what they produce.

' p Faith Integration Paper — Module 1

Assignment: Using Claude or Gemini as a co-writer, compose a 1-page reflection paper (approximately 300–400 words) responding to the following prompt:

You are about to become one of the most AI-capable business professionals in your field. Drawing on Genesis 2:15 and the concept of stewardship, describe the personal framework you will use to decide when, how, and for whose benefit you will deploy AI tools in your career. What guardrails will you put in place? Who are the "neighbors" your AI-powered work should serve?

Process:

• O
pe
n
Cla
ud
e
or
Ge
mi
ni
an
d s
har
e t
he
pro
mp

t a
bo
ve.

- Ask the AI to help you brainstorm, outline, and draft the paper.

- Revise the draft in your own voice —

add a personal story, conviction, or

ex
am
ple
th
at
ma
ke
s
it a
uth
ent
ical
ly y
our
s.

• S
ub
mit
th
e fi
nal
ve
rsi
on
(AI
-as
sist
ed,
hu
ma
n-o
wn
ed)
.

Format: 1 page, double-spaced, 12pt font, Times New Roman. Include the Genesis 2:15 passage at the top.

Grading: Assessed on theological engagement (does it genuinely wrestle with stewardship?), personal authenticity (is your voice present?), and practical specificity (does it name real decisions you will face?).

Lab 1: The Tool Gauntlet

Objective: Achieve baseline proficiency across all six core AI tools by completing a structured series of tasks that mirror real business use cases.

Estimated time: 3–4 hours (can be spread across multiple sessions)

Deliverable: A 500-word reflection document submitted to the course LMS describing your

experience, surprises, and initial mental model of each tool's strengths and limitations.

Lab Tasks

Task 1: Gemini — Market Landscape Analysis (45 min)

- Opening
emini.
googl
e.com
and
de
nab
le
De
ep
p
Research
arch
h
mo
de.
- Prompt: "Conduct a comprehensive analysis of the AI-powered red cu

sto
me
r s
erv
ice
m
ark
et
in
20
25
–2
02
6. I
de
ntif
y t
he
top
5
pla
yer
s, t
hei
r b
usi
ne
ss
mo
del
s,
pri
cin
g,
an
d k
ey
diff
ere
nti
ato
rs.
For
ma
t
as
a s
tru
ctu
red
re
por
t."

- Review the output. Identify why 2 claims you want to verify independently.
- Verify that those claims using NLP book LM (see Task 2) or direct web search.

Task 2: NotebookLM — Source Synthesis (30 min)

• U
plo
ad
3 d
oc
um
ent
s
to
not
eb
oo
kl
m.
go
ogl
e.c
om
:
(a)
an
y a
ca
de
mi
c p
ap
er
on
AI
in
bu
sin
es
s fr
om
G
oo
gle
Sc
hol
ar,
(b)
a c
om
pet
itor
's
an
nu
al r

ep
ort
or
10-
K fi
lin
g fr
om
S
EC
.go
v,
(c)
an
y n
ew
s a
rtic
le
ab
out

an
Al
bu
sin
es
s a
ppl
ica
tio
n fr
om
th
e p
ast
6
mo
nth
s.

- A
sk
No
teb
oo
kL
M:
"W
hat
ar
e t
he
thr

ee
mo
st i
mp
ort
ant
cl
ai
ms
th
es
e s
our
ce
s
ma
ke
ab
out
AI'
s i
mp
act

on
bu
sin
es
s?
Wh
ere

do
the
y a
gre
e a
nd
dis
agr
ee
?"

- N
ote
ho
w
No
teb
oo
kL
M
cit
es

so
urc
es.
C
om
par
e
to
ho
w
Ge
mi
ni r
es
po
nd
ed
in
Ta
sk
1.

Task 3: AI Studio — Prompt Engineering (30 min)

- O
pe
n a
ist
udi
o.g
oo
gle
.co
m
an
d s
ele
ct t
he
lat
est
G
em
ini
Pr
o
mo
del
.
- Wri
te
a p

rompt
to
analyze
a business
problem
of
your
choice.
Test
these
variations:
s:

-
A brief
, conversational
prompt
(1–2 sentences)

-
A structured
prompt

pt
wit
h r
ole
, c
ont
ext
, ta
sk,
an
d f
or
ma
t in
str
uct
ion
s

-
A c
hai
n-
of-t
ho
ug
ht
pro
mp
t th
at
as
ks
the
m
od
el
to r
ea
so
n s
tep

by
ste
p

- D
oc
um
ent
ho
w
out
put

quality changes across prompts.
This is your first prompt engineering experiment.

Task 4: Claude — Long-Context Analysis (30 min)

- Open Claude.ai and select the latest Claude model

.

- Upload this screenshot (or add document of 10,000+ words of your choice).

- Ask Claude: "Identify the three most contestable claims in this document"

. For each, provide the strongest suggestion to counter the argument, evidence that might support the claim, and your overall assessment of the claim's defensibility."

- E

na
ble
ex
ten
de
d t
hin
kin
g
if a
vail
abl
e.
Ob
ser
ve
ho
w t
he
rea
so
nin
g p
roc
es
s d
iffe
rs f
ro
m
sta
nd
ard
re
sp
on
se
s.

Task 5: Claude Code — First AI Coding Experience (45 min)

- In
sta
ll C
lau
de
Co
de
foll
owi
ng
the
se

tup
gu
ide
at
do
cs.
ant
hro
pic
.co
m/
cla
ud
e-c
od
e.

- C
rea
te
a n
ew
fol
der
: la
b0
1-d
ata
-an
aly
sis/

- P
ro
mp
t Cl
au
de
Co
de:
"C
rea
te
a P
yth
on
scr
ipt
tha
t d
ow
nlo
ad
s s

toc
k p
ric
e d
ata

for
thr
ee
co
mp
ani
es
of
my
ch
oic
e u
sin
g t
he
yfi
na
nc
e li
bra
ry,
cal
cul
ate
s t
hei
r
30-
da
y a
nd
90-
da
y
mo
vin
g a
ver
ag
es,
an
d g
en
era
tes
a
co

mp
ari
so
n c
har
t. I
ncl
ud
e c
om
me
nts
ex
pla
ini
ng
ea
ch
ste
p."
• R
un
the
ge
ner
ate
d c
od
e.
No
te
an
y e
rro
rs
an
d h
ow
Cl
au
de
Co
de
res
po
nd
s
to t
he
m.

Task 6: Tool Integration — Connecting the Stack (30 min)

- Take the market analysis from Task 1 (Gemini) and the sources synthesize from Task 2 (NotebookLM).
- Open Claude and paste both outputs.
- P
ro

mp
t:
"I h
av
e
an
AI
ma
rke
t a
nal
ysi
s a
nd
a s
ynt
he
sis
of
ac
ad
em
ic s
our
ce
s
on
AI
in
bu
sin
es
s. I
de
ntif
y t
he
thr
ee
big
ge
st
ga
ps
bet
we
en
wh
at i
nd
ust
ry
an

aly
sts
ar
e s
ayi
ng
an
d
wh
at
ac
ad
em
ic r
es
ear
ch
su
gg
est
s.
Wh
at
sh
oul
d
a b
usi
ne
ss
str
ate
gis
t p
ay
att
ent
ion
to
tha
t b
oth
so
urc
es
mi
ght

be
mi
ssi
ng
?"

- Reflect: what happened when you connected multiple AI tools in sequence? What was the output quality compared to any single tool?

AI Studio Build — Weekly Application

This is not optional. Every chapter in this course includes an AI Studio Build assignment. This is how you learn to build, not just analyze.

The BMC Disruption Analyzer

In aistudio.google.com, you will build a shareable AI application — your first of nine.

What you are building: A system-prompted Gemini application that functions as an industry-disruption analyst. When a user enters any industry name, the app returns a structured analysis of how AI disrupts each of the nine BMC blocks for that industry, with a specific real-world example per block.

Step-by-step:

- Open aistudio.google.com and navigate to "Build" (not "Stream Realtimed" or "Prompt").
- In the Sy

stem
instructional
field,
write
your
system
prompt. This
is
where
you
define
what
your
application
does,
how
it
behaves,
and
what
format
it
outputs.
(Ap

ply
m
eta
-pr
om
pti
ng
her
e

—
us
e
Cla
ud
e
to
hel
p y
ou
wri
te
a t
hor
ou
gh
sys
te
m
pro
mp
t b
efo
re
pa
sti
ng
it i
nto
AI
Stu
dio
)

•
A s
tro
ng
sys
te
m i
nst
ruc
tio
n

for
thi
s a
ppl
ica
tio
n
will
sp
eci
fy:
- T
he
AI'
s r
ole
(in
du
str
y d
isr
upt
ion
an
aly
st
wit
h d
ee
p b
usi
ne
ss
str
ate
gy
ex
per
tis
e)
- T
he
ex
act
ou
tpu
t st
ruc
tur
e (
on
e s
ect

ion
pe
r B
M
C
blo
ck,
wit
h
a c
on
cre
te
ex
am
ple

for
ea
ch)
- T
he
ton
e (
an
aly
tic
al,
evi
de
nc
e-b
as
ed,
sp
eci
fic
—

not
va
gu
e
or
ge
ner
ic)
- T
he
qu
alit
y b
ar (
ex

am
ple
s
mu
st
na
me
sp
eci
fic
co
mp
ani
es,
te
ch
nol
ogi
es,
or
bu
sin
es
s
mo
del
s
—
no
ab
str
act
ge
ner
aliti
es)
•T
est
yo
ur
ap
plic
ati
on
by
ent
eri
ng
at l
ea
st t
hre
e i

nd
ust
rie
s
of
yo
ur
ch
oo
sin
g.
Re
vie
w t
he
out
put
s c
riti
call
y:
Ar
e t
he
ex
am
ple
s s
pe
cifi
c?
Ar
e t
he
dis
rup
tio
n
me
ch
ani
sm
s a
cc
ura
te?
R
efi
ne
yo
ur
sys
te

m
pro
mp
t it
era
tiv
ely
un
til y
ou
are
sa
tisfi
ed.

• U
se
the
Sh
are
fu
nct
ion
in
AI
Stu
dio
to
ge
ner
ate
a
pu
blic
sh
are
lin
k.

Deliverable: Submit (a) your AI Studio share link, (b) the full text of your system instruction, and (c) sample outputs for three industries of your choosing. Your system instruction alone should be 2–3 substantial paragraphs.

What you are learning: System instructions, structured prompting, output format control, and iterative prompt refinement. These are the foundational skills for every AI Studio build in subsequent chapters.

Discussion Prompts

Discussion 1: The Winner's Curse in Your Industry

Select an industry you are familiar with (your current industry, a prior employer, or an industry you find strategically interesting). Apply the winner's curse framework from Section 1.1 to analyze the current AI adoption landscape in that industry.

In your post:

- Identify where your chosen industry currently sits

on the AI adoption curve (innovator, early adopter, early majority, late

majority

, or lagged)

- Cite at least one specific company that has moved early and describe the outcomes (positive or negative)

- Cite at least one

es
pe
cifi
c c
om
pa
ny
tha
t h
as
mo
ve
d l
ate
an
d d
es
cri
be
the
co
mp
etit
ive
co
ns
eq
ue
nc
es
• A
rgu
e
for
the
ad
opt
ion
str
ate
gy
yo
u b
eli
ev
e
is
opt
im
al
for
a n
ew

en
tra
nt
in t
his
in
du
str
y
in
20
26

Support your argument with scholarly sources on technology adoption and at least one recent news source (no more than two years old) documenting AI adoption patterns in your chosen industry.

Discussion 2: The BMC Disruption Audit

Select a company you admire or find strategically interesting. Conduct a focused BMC disruption audit.

In your post:

- Id
ent
ify
the
si
ngl
e
mo
st
dis
rup
ted
bl
oc
k
in
yo
ur
ch
os
en
co
mp
an
y's
BM
C
an
d e
xpl

ain
, wi
th
evi
de
nc
e,
wh
y y
ou
as
se
ss
it
as
the
m
ost
di
sru
pte
d

- Id
ent
ify
the
bl
oc
k
mo
st li
kel
y
to
be
co
me
a
co
mp
etit
ive
ad
va
nta
ge
if t
he
co
mp
an
y e
xe

cut
es
its
AI
str
ate
gy
wel
l

• A
rgu
e
wh
eth
er t
his
co
mp
an
y
is
mo
re l
ikel
y
to
be
an
AI
win
ner
or
los
er
ov
er t
he
ne
xt t
hre
e y
ear
s,
usi
ng
sp
eci
fic
evi
de
nc
e fr
om

public sources

Your post should demonstrate familiarity with at least one of these frameworks from your readings: Disruption theory (Christensen), Platform dynamics (Parker, Van Alstyne & Choudary), or Dynamic capabilities (Teece).

Readings

Required:

- Osterwalder, A., & Pigneur, Y. (2010). Business model generation: A handbook for visionaries, game

anger
ger
s,
an
d c
hal
len
ger
s.
Wil
ey.
(C
ha
pte
rs
1–
2:
Th
e
Ca
nv
as)
• C
hri
ste
ns
en,
C.
M.,
M
cD
on
ald
, R
, A
ltm
an,
E.
J.,
&
Pal
me
r, J
.E.
(20
18)
. Di
sru
pti
ve
inn
ov
ati

on:

An
int
ell
ect
ual
hi
sto
ry
an
d d
ire
cti
on
s
for
fut
ure
re
se
arc
h.

Jo
urn
al
of
Ma
na
ge
me
nt
Stu
die
s,
55(
7),
10
43
–1
07
8.

• M
cKi
ns
ey
Gl
ob
al I
nst
itut
e. (
20

25)

. The state of AI: How organizations are reengineering to capture value.

McKinsy & Company.

- EIU, London,

T., Manning,

S., Mishkin,

P., & Rock, D.

(20
23)
.G
PT
s a
re
GP
Ts:

An
ear
ly l
oo
k
at t
he
lab
or
ma
rke
t i
mp
act
po
ten
tial
of
lar
ge
lan
gu
ag
e
mo
del
s.
ar
Xiv
:23
03.
10
13
0.

Recommended:

- A
ndr
ee
ss
en,

M.

(20
11,
Au
gu
st
20)
. W
hy
sof
tw
are
is
eat
ing
th
e
wo
rld.
Th
e
Wa
ll S
tre
et
Jo
urn
al.
• B
ryn
jolf
ss
on,

E.,
&
Mc
Afe
e,
A. (
20
14)
. T
he
se
co
nd
ma
chi
ne
ag
e:
Wo
rk,

progress,
and prosper-
ity in a time
of brilliant
technologies.
W. W. Norton
& Company.
• Parker,
G. G., Van
Alstyne,
M. W., &
Choudhary,
S. P. (2016).
Platform

m r
ev
olu
tio
n:
Ho
w
net
wo
rke
d
ma
rke
ts
are
tra
nsf
or
mi
ng
the
ec
on
om
y a
nd
ho
w
to
ma
ke
the
m
wo
rk
for
yo
u.
W.
W.
No
rto
n
&
Co
mp
an
y.
• T
ee
ce,
D.
J. (

20
10)
. B
usi
ne
ss
mo
del
s,
bu
sin
es
s s
trat
eg
y a
nd
inn
ov
ati
on.
Lo
ng
Ra
ng
e P
lan
nin
g,
43(
2–
3),
17
2–
19
4.

Glossary

Term	Definition
Artificial Intelligence (AI)	The simulation of human intelligence processes by machines, particularly computer systems; in contemporary usage, refers primarily to machines that perform tasks that would otherwise require human intelligence, such as visual perception, speech recognition, decision-making, and translation.
Large Language Model (LLM)	A type of AI model, often based on transformer architecture, capable of generating human-like text across a wide range of topics and contexts.
Business Model Canvas (BMC)	A strategic management template developed by Alexander Osterwalder for developing and documenting business models, organized around nine essential elements.
Inflection Point	A point in time or a moment in a process where a significant change or trend occurs, often marking a turning point or a shift in direction.
Disruptive Innovation	Clayton Christensen's theory describing how new entrants with simpler, more affordable products can

displace established competitors by targeting overlooked segments before eventually moving upmarket

Segment of One	The theoretical limit of customer segmentation, in which each individual customer receives a fully
Adoption Curve	Everett Rogers' model describing how innovations spread through populations over time, with adopter
Winner's Curse	Based on the paradox where early movers face highest implementation costs and lowest model quality
Foundation Model	A large AI model trained on broad data that can be adapted to a wide range of downstream tasks;
Fine-Tuning	Examples include GPT-4o, Claude, and Gemini. Fine-tuning a foundation model on a domain-specific dataset to improve
Retrieval-Augmented Generation (RAG)	A technique combining AI generation with retrieval from external knowledge bases, enabling AI systems
Value Proposition	to address and benefit a specific customer segment by delivering tailored value. The answer to "why should a customer
Key Resource	In the BMC framework, the most important assets required to make a business model work — physical,
Cost Structure	the most important, or financial. The least important is the cost of operating under a particular business model
API (Application Programming Interface)	A set of protocols and tools that allows different software applications to communicate; in AI, the
Prompt Engineering	The practice of designing and optimizing prompts to elicit desired outputs from AI models; a key skill for
Marginal Cost	The cost of producing one additional unit of a good or service; AI is driving the marginal cost of knowledge
Moat (Competitive)	A company's sustainable competitive advantage that protects it from competitors; in the AI era, data assets
AI-Native	Describes a product, company, or workflow designed from the ground up with AI as a core component,
Outcome-Based Pricing	rather than a model where customers pay less on usage, achieved rather than access to a tool or platform; increasingly viable as AI makes outcomes
	more measurable and predictable