

BUSINESS APPLICATIONS OF AI

Module 2 — Value Proposition

Palm Beach Atlantic University · Graduate Program

Value Proposition

"A real opportunity is a problem that (1) everyone agrees is a problem, (2) can be solved at scale, and (3) you actually know how to solve. Without all three, you have a hobby."

Learning Objectives

By the end of this chapter, you will be able to:

- A ppl y t he "pa in-f irst " fr am ew ork to ide ntif y a nd arti cul ate cu sto me r pr obl em s w ith sp eci ficit y a nd em oti on al r es on

ance.

- Use the Three-Filter Test (Agreement, Scale, Competence) to evaluate and prioritize business opportunities.

- Map customer motivations using Job

s
to
Be
Do
ne
—
fun
cti
on
al,
em
oti
on
al,
an
d s
oci
al
—
an
d a
ppl
y t
he
"hir
ed
an
d fi
red
" in
ter
vie
w t
ec
hni
qu
e.
• D
isti
ng
uis
h tr
adi
tio
nal
co
mp
etit
ive
m
oat
s fr
om

Alternative models and evaluate the durability of each in a world where model capabilities come and go rapidly.

- Analyze the Cursor vs. GitHub Co

pilot case study to identify how workflow depth and positioning, not model access, drove differentiated business outcomes.

- Construct a complete V

al-
e P
rop
osi-
tio-
n
Ca-
nv-
as
for
a r-
eal
pr-
obl-
em
do-
ma-
in.

2.1 Pain First, Not Features First

Why 90% of Pitches Die in the First Two Minutes

Investor meetings are predictable. A founder sits down, opens with a slide deck, and within sixty seconds begins describing what their product does: "We built an AI-powered platform that integrates with your CRM to surface predictive insights and automate your outreach workflows." The room goes quiet. Not in anticipation — in polite disconnection.

The mistake is not technical. The pitch may describe a genuinely useful product. The mistake is sequencing: features before pain. When a founder leads with what they built, they are implicitly asking the audience to do the translation work — to imagine what problem this solves, for whom, and whether that problem is worth solving. Most audiences won't do that work. They haven't felt the pain, so they cannot evaluate the cure.

Studies of early-stage investor decision-making consistently find that the ability to articulate the problem clearly — not the sophistication of the solution — is the strongest early predictor of investment interest (Fried & Hisrich, 1994; Maxwell, Jeffrey, & Lévesque, 2011). The investor has heard hundreds of pitches about AI-powered platforms. They have not heard your specific user's pain articulated so precisely that it makes them flinch.

This is the "make them hurt, then heal them" arc — a structure borrowed from classical rhetoric but validated by every effective pitch deck study of the past decade. The sequence is:

```
flowchart LR
  A[Name the Pain] --> B[Make it Visceral]
  B --> C[Quantify the Cost]
  C --> D[Introduce the Cause]
  D --> E[Present the Solution]
  E --> F[Show the Moat]
```

problems we have felt than problems we have only been told about.

Paint the Pain: The Sarah Exercise

The single most powerful tool for a first-time founder is what we call the Sarah Exercise. It is deceptively simple: replace every abstract, category-level pain description with a named individual experiencing a specific moment of acute suffering.

Here is the exercise in action.

Abstract pain description (typical founder language):

"Small and mid-size retailers face significant inventory management challenges, leading to inefficiencies and lost revenue opportunities."

This sentence is technically true of approximately 500,000 businesses in the United States. It is also completely inert. No investor or customer has ever been moved by the phrase "inventory management challenges."

Sarah version:

"It's 2:17 in the morning on the Tuesday before Thanksgiving. Sarah Chen, who owns a kitchenware store in Columbus, Ohio, wakes up to find an email from her largest wholesale customer — a regional restaurant chain — ordering \$50,000 worth of product for the holiday season. Sarah opens her inventory system and discovers she has 40% of the required stock. She cannot fulfill the order. She emails back to apologize. She loses the order. She loses the customer. She lies awake for the next three hours calculating how she'll make payroll in January."

The abstract version describes a category. The Sarah version describes a human being at the worst moment in their professional week. The difference in persuasive force is not small — it is categorical.

The Sarah Exercise works because it forces you to answer several questions you should have answered before writing a single line of code:

-

Wh
o s
pe
cifi
call
y
is
ex
per
ien
cin

g t
his
pa
in?
N
ot
a d
em
ogr
ap
hic
se
gm
ent

—
a p
ers
on
wit
h
a n
am
e,
a r
ole
, a
nd
a b
ad
Tu
es
da
y.

•
Wh
en
ex
act
ly
do
es
the
pa
in
oc
cur
? P
ain
ha
s
a
mo

ment.
That
at
moment
is your
product
's
highest-
level
average
intervention
point.

- What
does
it cost
? Not
necessarily
—
specifically
\$.50,
000
orders,
three
hours

of
lost
sleep,
once
customer
relationships
hip
destroyed.

- What
is the
emotional
consequence?
Sarah
isn't
just
experiencing
a business
setback.
She's

ex
per
ien
cin
g f
ear
, s
ha
me
, a
nd
the
pa
rtic
ula
r dr
ea
d
of
sm
all-
bu
sin
es
s o
wn
ers
hip
.

Pain Quantification: The Four Dimensions

Once you have your Sarah, you need to measure her pain. Investor conversations and customer discovery interviews require concrete numbers, not impressions. Pain quantification operates across four dimensions:

Dimension	Measure	Example
Time wasted	Hours/week per affected person	Sarah spends 12 hours/week on manual inventory reconciliation
Dollars lost	Revenue lost, costs incurred	\$50,000 order lost; \$3,400/month in excess carrying costs
Deals missed	Opportunities not captured	2–3 large orders declined per quarter due to stock uncertainty
Sleep lost	A proxy for emotional cost	3 hours of anxiety-sleep after inventory failures

The "sleep lost" dimension is not a joke. It is a shorthand for the emotional weight of a problem — the degree to which it occupies mental space outside work hours. Problems that cost money but don't keep anyone up at night are commodity inconveniences; someone will solve them with a \$9/month SaaS tool. Problems that keep people up at night are the kind that justify premium pricing, long-term contracts, and deep customer loyalty.

2.2 The Three-Filter Test

Most aspiring founders have more ideas than discipline. The Three-Filter Test is a structured methodology for killing ideas quickly — before they consume months of your life. Its premise is that a genuine business opportunity must satisfy three conditions simultaneously. Failing any one of them is disqualifying.

Filter 1: Agreement

The question: Do ten unrelated people name the same pain in the same words without being prompted?

This filter is far more demanding than it appears. Many founders conduct "customer discovery" interviews and receive polite agreement: "Yes, that is a real problem." Polite agreement is not the same as unprompted convergence. The agreement filter requires that when you ask ten people in your target market to describe their biggest operational challenge, a significant proportion of them independently reach for the same vocabulary.

When Rob Fitzpatrick, author of *The Mom Test* (2013), developed his framework for evaluating customer feedback, he observed that most customer discovery is contaminated by what he called the "mom problem" — the tendency of potential customers to tell founders what they want to hear rather than what is true. The Agreement filter bypasses this by looking for convergence, not affirmation.

How to run it:

- Identify 10 potential customers in your target market. Do not pitch your i

de
a.

• A
sk
op
en-
en
de
d q
ue
sti
on
s: "
Wh
at'
s t
he
har
de
st
par
t
of
yo
ur
we
ek
?"

"W
her
e
do
yo
u l
os
e t
he
mo
st t
im
e?"
"

Wh
at
pro
ble
m
wo
uld
yo
u p
ay
\$1,

00
0
to
sol
ve
rig
ht
no
w?
"

- R
ec
ord
th
e e
xa
ct
wo
rds
th
ey
us
e.
Do
no
t p
ara
phr
as
e.

- L
oo
k
for
co
nv
erg
en
ce:

do
thr
ee
or
mo
re
pe
opl
e d
es
cri
be

the
sa
me
pa
in
usi
ng
si
mil
ar l
an
gu
ag
e?

If the top answer is scattered — everyone has a different worst problem — the market is not ready for a point solution. If the top answer converges — "we can never get accurate inventory counts" appears in 6 of 10 conversations — you have passed Filter 1.

Filter 2: Scale

The question: Does this problem exist for at least 10,000 addressable entities who could plausibly pay for a solution?

The 10,000 threshold is a practical minimum for a venture-scale business. A problem that affects 200 businesses is a consulting engagement. A problem that affects 10,000 businesses — each paying \$2,400/year — is a \$24 million ARR business at full penetration, which, even at 10% market share, is a \$2.4 million ARR business. That is not a transformative enterprise, but it is a real business.

The scale filter requires a bottom-up market sizing exercise:

```
flowchart TD
    A["Total Addressable Market (TAM)"] --> B["Who has this exact problem?"]
    B --> C["Serviceable Addressable Market (SAM)"] --> D["Who can we reach given our GTM?"]
    D --> E["Serviceable Obtainable Market (SOM)"] --> F["What's realistic in years 1-3?"]
```

Worked example: Suppose you want to build an AI-powered scheduling tool for independent physical therapy practices. Your bottom-up sizing might look like:

• T
AM
: 3
9,0
00
ind
ep
en
de
nt
ph
ysi

cal
th
era
py
pra
ctic
es
in t
he
US
(B
ure
au
of
La
bor
St
ati
stic
s d
ata
)

• S
AM
: 2
1,0
00
pra
ctic
es
wit
h
2+
the
rap
ists
(s
ma
ller
pr
act
ice
s w
ill n
ot
pa
y
for
sof
tw
are
)

• S
O
M:
2,1
00
pra
ctic
es
(10
%
ma
rke
t s
har
e
in
ye
ars
1–
3,
giv
en
dir
ect
sa
les
+ r
efe
rral
m
oti
on)
• P
ric
e p
oin
t: \$
29
9/
mo
nth
=
\$3,
58
8/y
ear
• S
O
M r
ev
en
ue:

2,
10
0
× \$
3,5
88
= ~
\$7.
5M
A
RR

\$7.5M ARR is not a unicorn. But it is a real business, and if the problem passes Filter 1 (physical therapists consistently naming scheduling as their biggest operational pain), it passes Filter 2.

Filter 3: Competence

The question: Do you have unfair advantages in solving this problem?

The word "unfair" is deliberate. A business advantage that any well-funded team could replicate in six months is not a moat — it is a head start, and head starts erode. Unfair advantages are asymmetric: they derive from who you are, what you know, who you know, or what data you have access to, not just what you built.

Competence can take several forms:

Domain expertise You have spent years inside the problem domain. You know the vocabulary, the workflows, the regulatory environment, and the specific failure modes. Competitors would need to hire 10 domain experts to replicate your intuition. Example: A former hospital administrator building AI tools for hospital billing has an unfair advantage over a generic fintech team.

Running the Three-Filter Test: A Worked Example

Let us apply all three filters to a candidate idea: an AI assistant for freelance graphic designers to automate client briefing and revision management.

Filter 1 — Agreement: We interview 12 freelance graphic designers and ask: "What's the most frustrating part of client work?" Responses:

• 9/
12
me
nti
on
so
me
ve
rsi
on
of "
un
cle

ar
bri
efs
"

or "
co
nst
ant
re
visi
on
req
ue
sts
"

• 7/
12
us
e t
he
phr
as
e "
sc
op
e c
ree
p"

• 5/
12
sp
eci
fic
ally
m
ent
ion
w
ast
ing
ho
urs
re
sp
on
din
g
to
em
ails
as
kin

g
for
the
sa
me
inf
or
ma
tio
n t
he
y a
lre
ad
y p
rov
ide
d

Verdict: Passes Filter 1. The pain is converging around unclear briefs, scope creep, and communication overhead.

Filter 2 — Scale:

- F
ree
lan
ce
gra
phi
c d
esi
gn
ers
in
the
U
S:
~3
30,
00
0 (
Bu
rea
u
of
La
bor
St
ati
stic
s)
- D

esi
gn
ers
do
ing
pri
ma
rily
cli
ent
w
ork
(v
s.
in-
ho
us
e):
~1
80,
00
0

• D
esi
gn
ers
ea
rni
ng
\$4
0,0
00
+/y
ear
(c
an
aff
ord
so
ftw
are
): ~
95,
00
0

• P
ric
e p
oin
t: \$
49/
mo

nth
 =
 \$5
 88/
 ye
 ar
 •
 At
 5%
 pe
 net
 rati
 on:
 4,
 75
 0 c
 ust
 om
 ers
 ×
 \$5
 88
 = ~
 \$2.
 8M
 A
 RR

Verdict: Marginal. \$2.8M ARR is a small business, not a venture opportunity. However, the total addressable market expands significantly if the tool is designed for all freelance creative professionals (photographers, copywriters, web designers) — expanding the pool to potentially 500,000+ professionals. With that framing, it passes Filter 2.

Filter 3 — Competence:

• T
 ea
 m
 me
 mb
 er
 A:
 For
 me
 r cr
 eat
 ive
 dir
 ect
 or
 at
 a
 mi

d-s
ize
ag
en
cy
wit
h
12
ye
ars
of
clie
nt-f
aci
ng
cre
ati
ve
wo
rk

- T
ea
m
me
mb
er
B:
Ful
l-st
ac
k d
ev
elo
per
wit
h
UX
ba
ck
gro
un
d

- D
om
ain
ex
per
tis
e:
Str
on
g (f

or
me
r C
D
ha
s li
ve
d t
he
pro
ble
m)

- P
rop
riet
ary
da
ta:
No
ne
yet
, b
ut
pai
d p
ilot
ac
ce
ss
to
3 a
ge
nci
es
for
dat
a c
oll
ect
ion

- D
istr
ibu
tio
n:
Te
am
A
ha
s L
ink
edl

n n
et
wo
rk
of
2,4
00
de
sig
ner
s a
nd
cre
ati
ve
dir
ect
ors

Verdict: Passes Filter 3. Domain expertise is genuine, distribution is above average for an early team.

Outcome: This idea passes all three filters. It is not a billion-dollar idea in its current framing, but it is a real opportunity. The team should proceed to customer validation.

2.3 Jobs to Be Done

Christensen's Framework in the AI Era

Clayton Christensen introduced the Jobs to Be Done (JTBD) framework in *The Innovator's Solution* (2003) and expanded it in *Competing Against Luck* (2016). The central insight is disarmingly simple: customers don't buy products — they hire outcomes.

When a customer "hires" a product, they are hiring it to make progress on a specific job in their life or work. When they "fire" it, they are firing it because it failed to deliver that progress, or because a better-suited solution became available. The product is a means to an end. The job is the end.

This reframing has profound implications for how you design and pitch a value proposition. A drill manufacturer who understands that customers are "hiring a drill" will compete on drill specifications: RPM, battery life, chuck size. A drill manufacturer who understands that customers are "hiring a hole in the wall" will ask different questions: Why do they need the hole? Could an adhesive hook serve the same job? Could a professional installation service? The JTBD lens constantly pushes the question from "what does our product do?" to "what progress is the customer trying to make?"

In the AI era, the JTBD framework has taken on new relevance. AI tools are generative — they can complete a very wide range of tasks. This creates a temptation to define the job as "use AI to do X." But "use AI to do X" is a feature, not a job. The job is always upstream of the feature. A legal research AI does not exist so lawyers can "use AI for research" — it exists so lawyers can deliver confident, well-researched legal opinions to clients faster and with less anxiety.

Three Dimensions of Jobs

Every customer job has three dimensions:

Functional Jobs The practical, task-level outcome the customer needs to achieve. These are the most visible and the most commonly addressed in product design — but they are also the most easily replicated by competitors.

Example: Sarah needs to know her current inventory levels across all SKUs, updated in real time, accessible from her phone.

Emotional Jobs How the customer wants to feel after hiring your product. These are rarely stated in customer interviews but are consistently present in purchase decisions. Emotional jobs are where differentiation lives — two products with identical functional capabilities will be separated by which one makes the customer feel more capable, more in control, less anxious.

Example: Sarah wants to feel like a competent, organized business owner — not like someone perpetually scrambling to catch up with her own inventory.

Social Jobs How the customer wants to be perceived by others — customers, employees, peers, investors. Social jobs are particularly important in B2B contexts, where purchase decisions often involve organizational signaling as much as functional need.

Example: Sarah wants to be the kind of business owner that restaurant chain buyers trust to deliver reliably. The inventory tool is, in part, a signal of professionalism.

The "Hired and Fired" Interview Technique

The most powerful method for uncovering jobs is the Hired and Fired Interview, a structured qualitative technique developed from Christensen's work and popularized by Bob Moesta and Chris Spiek. The technique focuses on two types of switching events:

- Hired and Fired Interview: Tell me about the last time you started using a new

to
ol
or
ser
vic
e
to
sol
ve
thi
s p
rob
le
m.
Wa
lk
me
thr
ou
gh
ex
act
ly
wh
at
ha
pp
en
ed
in t
he
da
ys
an
d
we
ek
s b
efo
re
yo
u
ma
de
tha
t d
eci
sio
n.
• Fi
red
: T
ell

me
ab
out
th
e l
ast
ti
me
yo
u s
top
pe
d u
sin
g
a t
ool
or
ser
vic
e y
ou
ha
d b
ee
n u
sin
g
for
thi
s p
rob
le
m.
Wh
at
ha
pp
en
ed
?

The goal is to reconstruct the purchase story and the abandonment story in granular detail.
Purchase stories reveal:

- T
he
trig
ger
ing
ev
ent
(w

that
changed
that
made
the
problem
urgent
enough
to
act

on)

- The struggling moment (the specific instance of pain that motivated the search)

- The

consideration set (what alternative they evaluated and why they reject them)

- The hiring criteria (what made them ultimately choose this solution)

Abandonment stories reveal:

- The failure mode (what he produced stopped people delivering)

- The competing hire (what he switched to and why)

- The unmet emotional job (what

at t
he
y f
elt
the
pr
od
uct
ne
ver
re
ally
ga
ve
the
m)

Why Customers Don't Buy Products — They Hire Outcomes

The implications of JTBD for value proposition design are direct:

- Y
our
va
lue
pr
op
osi
tio
n s
ho
uld
na
me
th
e j
ob,
no
t th
e p
rod
uct
. In
ste
ad
of "
inv
ent
ory
m
an
ag

em
ent
so
ftw
are

for
ret
ail
ers
," c
on
sid
er "
sle
ep
ins
ura
nc
e
for
sm
all
bu
sin
es
s o
wn
ers
w
ho'
ve
be
en
bur
ne
d
by
sto
cko
uts
." T
he
se
co
nd
fra
mi
ng
na
me
s t
he

em
oti
on
al j
ob
(sl
ee
p
—
saf
ety

—
co
ntr
ol)
an
d t
he
trig
ger
ing
ex
per
ien
ce
(be
ing
bu
rne
d)
whi
le i
mp
lyin
g t
he
fun
cti
on
al
ca
pa
bilit
y.

• Y
our
co
mp
etit
or
set
is

bro
ad
er t
ha
n y
our
pr
od
uct
ca
teg
ory
. If
the
jo
b
is "
fee
l
in
co
ntr
ol
of
my
bu
sin
es
s,"
yo
u a
re
not
co
mp
eti
ng
onl
y w
ith
oth
er i
nv
ent
ory
so
ftw
are

—
yo
u a
re

competing with hiring an operations manager, using spreadsheets so businesses take management courses, or simply accepting the anxiety as

a c
ost
of
ow
ner
shi
p. J
TB
D r
ev
eal
s t
he
full
co
mp
etit
ive
la
nd
sc
ap
e.

• T
he
em
oti
on
al j
ob
det
er
mi
ne
s t
he
pri
ce
ceil
ing
. C
ust
om
ers
wil
l p
ay
mo
re
to
eli
mi
nat

e a
nxi
ety
th
an
to i
mp
rov
e e
ffici
en
cy.
If y
our
pr
od
uct
so
lve
s t
he
fun
cti
on
al j
ob
(sa
ve
s
12
ho
urs
/w
ee
k),
the
cu
sto
me
r c
alc
ula
tes
th
e p
ric
e
of t
hei
r ti
me
an
d s

ets
a
pri
ce
ceil
ing
ac
cor
din
gly.
If
yo
ur
pro
du
ct
sol
ve
s t
he
em
oti
on
al j
ob
(eli
mi
nat
es
the
2a
m
pa
nic
) , t
he
pri
ce
ceil
ing
is
mu
ch
hig
her

—
be
ca
us
e t
her
e

is
no
ho
url
y r
ate

for
pe
ac
e
of
mi
nd.

2.4 Moats in an AI-Native World

Traditional Moats and AI-Era Durability

In classical competitive strategy, a moat is a durable structural advantage that allows a company to generate above-market returns for an extended period. Warren Buffett popularized the term; Michael Porter's Five Forces framework provides the analytical scaffolding. Traditional moats include:

Moat Type	Traditional Strength	AI-Era Durability	Why
Brand	Very high	Moderate	AI can replicate brand voice and aesthetic; AI does not dissolve network effects as in fact
Network effects	Very high	High	AI-native workflows create *deeper* switching costs than legacy ones
Switching costs	High	High	AI accelerates reverse engineering; patent
Scale economies	High	Moderate–Low	AI improves content production but does not inherently open distribution channels
Proprietary IP / patents	Moderate	Low	AI improves content production but does not inherently open distribution channels
Proprietary data	Moderate	Very high	AI improves content production but does not inherently open distribution channels
Regulatory capture	Moderate	High	AI improves content production but does not inherently open distribution channels
Exclusive distribution	Moderate	Moderate	AI improves content production but does not inherently open distribution channels

The table reveals a clear pattern: moats that derive from structural market position (networks, switching costs, regulatory) hold in the AI era. Moats that derive from knowledge production efficiency (scale, IP) erode — because AI democratizes knowledge production.

New AI-Native Moats

The AI era has also created a new class of moat with no direct equivalent in prior competitive strategy frameworks. These AI-native moats are worth understanding in depth because they are poorly recognized by traditional investors and frequently overlooked by founders who focus too

heavily on model capability.

Moat 1: Workflow Depth

Workflow depth describes how completely a product has been woven into the daily operational fabric of the customer's work. A tool that an employee opens once a week to generate a report is a peripheral product — replaceable with limited disruption. A tool that an employee interacts with 47 times per day, at every decision point, that has their historical context, their preferences, their team's norms embedded in its outputs — that tool is operationally irreplaceable.

The depth of workflow integration is the primary determinant of switching costs in AI-native products. This is why Notion's AI features are defensible even against more capable standalone AI tools: they are embedded in the document where the work lives. This is why GitHub Copilot, despite its first-mover advantage, was vulnerable to Cursor — as we will explore in the case study below.

Moat 2: Proprietary Feedback Loops

Every interaction a user has with your AI product is a training signal. Every correction, every re-generation, every accepted or rejected suggestion teaches the system something about what good outputs look like for this specific user, company, or domain. Over time, these feedback signals accumulate into a proprietary preference dataset that no competitor can acquire — because it was generated by your users, doing their work, in your product.

This is categorically different from training data that any team can purchase or scrape from the internet. Proprietary feedback loops create personalization depth that compounds over time. A system that has processed 50,000 of Sarah's inventory decisions knows things about how Sarah evaluates tradeoffs that no general model knows — and no competitor can acquire that knowledge without replicating Sarah's history.

Duolingo is a clear example of this moat at scale. Their AI-powered language instruction engine has been trained on hundreds of millions of learning interactions — what spacing intervals work for which types of learners, which error patterns predict future failure, which exercise types produce retention. That dataset cannot be purchased. It can only be generated by having 40 million active learners using the product daily.

Moat 3: Agent Graph Position

As AI systems evolve from single-model tools to multi-agent architectures — networks of specialized AI agents that coordinate to complete complex tasks — position in the agent graph becomes a competitive advantage. An agent that coordinates other agents, routes tasks, and manages state across a workflow holds a structurally superior position to a specialist agent that only performs one function.

This is analogous to platform economics: the entity that controls the coordination layer captures more value than the entities that provide specific capabilities. A company whose AI system sits at the "orchestration layer" of a customer's agent workflow — routing tasks to specialized agents, maintaining memory, managing the human-in-the-loop interface — has a position that is extremely difficult to displace.

This moat is still emerging. As of 2025, most enterprise AI deployments are single-model or loosely coupled tools. The companies that establish orchestration-layer positions in the next 24 months will be extremely difficult to dislodge once multi-agent workflows become standard operating procedure.

Moat 4: Prompt IP and Evals

The final AI-native moat is the least glamorous but arguably the most undervalued: accumulated prompt infrastructure and evaluation capability. This includes:

- P

rompt
libraries:
Curated,
tested
systems
prompt
that
are
reliably
produce
high-
quality
outputs

for
specific
tasks

- E

valuation:
Automated
test
suite

hat
ev
alu
ate
w
het
her
Al
out
put
s
me
et
qu
alit
y s
tan
dar
ds

—
en
abl
ing
ra
pid
m
od
el
up
gra
de
s w
ith
out
qu
alit
y r
egr
es
sio
n

• Fi
ne-
tun
ing
da
tas
ets
: L
ab
ele
d d
ata

set
s d
eri
ve
d fr
om
hu
ma
n-r
evi
ew
ed
out
put
s t
hat
ca
n
be
us
ed
to f
ine
-tu
ne
mo
del
s
for
sp
eci
fic
do
ma
ins
• F
ail
ure
m
od
e d
oc
um
ent
ati
on:
In
stit
uti
on
al
kn
owl

ed
ge
of
wh
ere
th
e
mo
del
fai
ls,
un
der
w
hat
co
ndi
tio
ns,
an
d
wh
at
gu
ard
rail
s p
rev
ent
th
os
e f
ail
ure
s

A team that has spent 18 months building, testing, and refining prompt infrastructure for clinical documentation review has an asset that a new competitor starting from scratch cannot replicate in 6 months — even if they access the same underlying model. The prompt infrastructure represents accumulated knowledge about how to make the model work reliably in a specific high-stakes domain.

Case Study: Cursor vs. GitHub Copilot

In 2023, GitHub Copilot was the dominant AI coding assistant. It had first-mover advantage, the backing of Microsoft and OpenAI, and integration into the world's most widely used code editor (VS Code). It seemed unassailable.

By 2025, Cursor had captured a significant and vocal share of professional developers — particularly in startup environments — and was generating substantial ARR with a fraction of

Copilot's institutional resources. Both products were using similar underlying models. How?

The Setup: Same Models, Different Architectures

GitHub Copilot was designed as a code completion plugin — it operated inside VS Code, suggesting completions as the developer typed. The mental model was "smart autocomplete." The job being hired was narrow: finish my line of code faster.

Cursor was designed as a programmer's AI partner — it was itself an IDE (based on VS Code's open-source core), meaning the AI was not a plugin inside someone else's product but the environment itself. The mental model was "AI-native development." The job being hired was broader: help me build software faster, from architecture through debugging.

The Moat Analysis

Cursor's moats:

-

Wo
rkfl
ow
de
pth
: B
ec
au
se
Cu
rso
r
is t
he
ID
E r
ath
er t
ha
n
a p
lug
in i
nsi
de
on
e,
it h
as
ac
ce
ss
to t
he

full
pr
oje
ct
co
nte
xt
—
ev
ery
fil
e,
ev
ery
fu
nct
ion
, e
ver
y d
ep
en
de
nc
y.
Co
pil
ot,
op
era
tin
g
as
a p
lug
in,
ha
d li
mit
ed
co
nte
xt
win
do
ws
tha
t re
stri
cte
d h
ow
m

uc
h
of t
he
co
de
ba
se
it c
oul
d "
se
e."
Cu
rso
r's
Al
co
uld
re
sp
on
d
to i
nst
ruc
tio
ns
like
"re
fac
tor
the
au
the
nti
cat
ion
m
od
ule
to
us
e J
W
T i
nst
ea
d
of
se
ssi
on

co
oki
es
acr
os
s t
he
ent
ire
co
de
ba
se"

—
a t
as
k r
eq
uiri
ng
full
-pr
oje
ct
visi
bilit
y t
hat
C
opi
lot
co
uld
no
t re
lia
bly
co
mp
let
e.

- P
rop
riet
ary
fe
ed
ba
ck l
oo
ps:
Ev

ery
int
era
cti
on
Cu
rso
r u
ser
s h
av
e

—
ev
ery
ac
ce
pte
d s
ug
ge
sti
on,
ev
ery
rej
ect
ion
, e
ver
y
ma
nu
al
edi
t

—
ge
ner
ate
s tr
ain
ing
si
gn
al t
hat
im
pro
ve
s C
urs
or'

s
mo
del

for
su
bs
eq
ue
nt
us
ers

. O
ver
ti
me

, th
is f
ee
db
ac
k a
cc
um
ula
tes

int
o p
rop
riet
ary
fin
e-t
uni
ng
dat
a

for
co
de
pat
ter

ns,
err
or
cor
rec
tio
ns,

an
d d
ev
elo

per
pr
efe
ren
ce
s.

• A
ge
nt
gra
ph
po
siti
on:
C
urs
or'
s "
Co
mp
os
er"
fea
tur
e (l
ate
r e
xp
an
de
d
as
"A
ge
nt"
mo
de)
po
siti
on
s t
he
AI
at t
he
tas
k-o
rch
est
rati
on
lev
el:

it can create files, run the minimal commands, read the test outputs, and iterate on code automatically. This is not a code completion plugin

—
it is a dev

elo
pm
ent
w
ork
flo
w
orc
he
str
ato
r. T
hat
or
ch
est
rati
on
po
siti
on
is
ext
re
me
ly
diff
icul
t
for
a p
lug
in-
ba
se
d c
om
pet
itor
to
rep
lica
te
wit
ho
ut r
eb
uil
din
g t
he
pro
du

ct f
ro
m
scr
atc
h.

- S
wit
chi
ng
co
sts
: D
ev
elo
per
s w
ho
ha
ve
us
ed
Cu
rso
r
for
3+
mo
nth
s h
av
e t
hei
r e
dit
or
pre
fer
en
ce
s,
key
bin
din
gs,
an
d
wo
rkfl
ow
pa
tter
ns

em
be
dd
ed
in t
he
too
l. T
he
swi
tch
ing
co
st
to r
etu
rn
to
VS
C
od
e
+
Co
pil
ot
is
no
n-tr
ivia
l
—
not
be
ca
us
e
of
dat
a l
oc
k-
in,
but
be
ca
us
e
of
wo
rkfl
ow
ha

bit
loc
k-
in.

What Copilot missed:

GitHub Copilot's design reflected a conservative product philosophy: do not disrupt the existing workflow, add AI incrementally to what developers already use. This is a sensible risk-minimization strategy that works well in stable markets. In a rapidly evolving AI market, it produced a product that was incrementally better within a workflow architecture that was structurally limiting.

Copilot hired itself as a "typing assistant." Cursor hired itself as a "software development partner." The jobs are not the same, and the products that emerge from those job definitions are not the same.

What the Moat Was

The Cursor moat was not a better AI model. It was:

- Full-project context (technical architecture choice enabling workflow depth)
- Orchestrated on-lay

er
po
siti
oni
ng
(ag
ent
gr
ap
h p
osi
tio
n)

- A
cc
um
ula
ted
de
vel
op
er f
ee
db
ac
k (
pro
pri
eta
ry f
ee
db
ac
k l
oo
ps)

- Wo
rkfl
ow
ha
bit
for
ma
tio
n (
swi
tch
ing
co
sts
)

This is the template for AI-native competitive advantage. The companies that win in the AI era will not win because they have access to better models — models are increasingly available to everyone. They will win because they have built systems that accumulate proprietary advantage the longer customers use them.

Faith Integration — Module 2

*"Whatever you do, work at it with all your heart, as working for the Lord, not for human masters."
— Colossians 3:23 (NIV)*

Excellence as Worship: The Theology of a Great Value Proposition

Colossians 3:23 is one of the most practically disruptive verses in the New Testament. It doesn't say "do your best when it matters." It says whatever you do — including, implicitly, the spreadsheets, the pitch decks, the customer interviews, and the product specs — do it wholeheartedly, as if the audience is God himself.

This verse reframes what it means to build a great value proposition. A pain-first value proposition is not merely good strategy; it is an act of attention — of actually listening to another person's struggle rather than projecting your solution onto them. The Jobs-to-Be-Done framework is, at its core, an exercise in empathy. What is this person really trying to accomplish? What is getting in their way? What would it mean for their life if that problem disappeared?

Christian business professionals have a theological basis for this kind of deep customer attention that secular frameworks cannot fully supply. We are commanded to love our neighbors as ourselves (Matthew 22:39). Your customer is your neighbor. Building a value proposition that genuinely solves their problem — not one that merely generates revenue — is an act of love expressed through commerce.

The standard Colossians sets is demanding precisely because it refuses to let us compartmentalize: our professional craft is not separate from our faith. The quality of care we put into our value proposition is a form of worship.

'b Faith Integration Paper — Module 2

Assignment: Using Claude or Gemini as a co-writer, compose a 1-page reflection paper (approximately 300–400 words) responding to the following prompt:

Colossians 3:23 calls you to work wholeheartedly "as for the Lord." Apply this to the value proposition you are developing in this course. What would it look like to build a product or service for your target customer with that same standard of care? Where are you tempted to cut corners on genuinely understanding their pain? What would change if you treated customer discovery as a spiritual discipline?

Process:

- S
har

e t
he
pro
mp
t wi
th
Cla
ud
e
or
Ge
mi
ni
an
d u
se
the
AI
to
bra
ins
tor
m
an
d d
raft

• R
evi
se
int
o y
our
o
wn
vo
ice

— i
ncl
ud
e
a r
eal
ex
am
ple
fro
m
yo
ur
cu
sto

me
r re
se
arc
h
or l
ife
ex
per
ien
ce.
• S
ub
mit
th
e fi
nal
, h
um
an-
ow
ne
d v
ers
ion
.

Format: 1 page, double-spaced, 12pt font, Times New Roman. Include Colossians 3:23 at the top.

Grading: Theological engagement, personal authenticity, connection to your actual value proposition work.

Lab 2: The Value Proposition Canvas

Estimated time: 3–4 hours

Overview: The Value Proposition Canvas (VPC) is a tool developed by Alexander Osterwalder and Yves Pigneur that maps customer jobs, pains, and gains against a product's value map. It is the most widely used structured framework for designing and evaluating value propositions in the startup ecosystem.

In this lab, you will complete a VPC for a real problem domain, write a one-paragraph value proposition, and conduct a preliminary moat analysis.

Deliverable: (1) A completed Value Proposition Canvas (template or hand-drawn, photographed), (2) a one-paragraph value proposition statement, and (3) a moat analysis identifying which of the four AI-native moats your proposed solution could develop.

Lab Tasks

Task 1: The Sarah Exercise — Persona and Pain Narrative (45 min)

- Select a business domain you know well — your industry, a prior employer's industry, or a market you have researched.
- Identify the

single biggest operational national pipeline in that domain.
Write it as a category-level description on first (one sentence).

- Transform it using the Sarah Exercise: write a

15
0—
20
0
wo
rd
nar
rati
ve
na
mi
ng
a s
pe
cifi
c p
ers
on,
at
a s
pe
cifi
c
mo
me
nt,
ex
per
ien
cin
g a
cut
e p
ain
. In
clu
de
at l
ea
st
on
e q
ua
ntif
ied
co
st.
• U
se
the
fol
low
ing

st
art
er
pro
mp
t wi
th
yo
ur
Al t
ool
of
ch
oic
e (
aft
er
me
ta-
pro
mp
tin
g
it):
"H
elp
m
e tr
an
sfo
rm
thi
s c
ate
gor
y-l
ev
el
pai
n d
es
cri
pti
on
int
o
a v
ivid
, s
pe
cifi
c p
ers

on
a n
arr
ati
ve:
[p
ast
e y
our
de
scr
ipti
on]
. T
he
nar
rati
ve
sh
oul
d n
am
e
a s
pe
cifi
c p
ers
on,
de
scr
ibe
a
sp
eci
fic
ba
d
mo
me
nt,
incl
ud
e
a q
ua
ntif
ied
co
st,
an
d s
urf

ac
e t
he
em
oti
on
al
co
ns
eq
ue
nc
e."

Task 2: Customer Discovery — The Three-Filter Test (60 min)

- T
ak
e t
hre
e c
an
did
ate
bu
sin
es
s i
de
as
fro
m
yo
ur
ow
n e
xp
eri
en
ce,
re
se
arc
h,
or i
nte
res
t.
- F
or
ea
ch
ide

a,
co
mp
let
e t
he
Thr
ee-
Filt
er
Te
st:
- Fi
lter
1 (
Ag
ree
me
nt):
W
hat
ev
ide
nc
e
do
yo
u h
av
e t
hat

10
pe
opl
e n
am
e t
his
pa
in
un
pro
mp
ted
?
If y
ou
lac
k e
vid
en
ce,

de
sig
n
a
5-q
ue
sti
on
int
erv
iew
pr
oto
col
to
gat
her
it.
- Fi
lter
2 (
Sc
ale
):
Co
nd
uct
a
bot
to
m-
up
ma
rke
t si
zin
g u
sin
g a
vail
abl
e p
ubl
ic
dat
a.
Sh
ow
yo
ur
ma
th.
- Fi

lter
3 (
Co
mp
ete
nc
e):
Lis
t y
our
te
am
's
unf
air
ad
va
nta
ge
s.
Be
ho
ne
st
ab
out
ga
ps.
• U
se
thi
s s
tart
er
pro
mp
t (a
fter
m
eta
-pr
om
pti
ng)
: "I'
m
ev
alu
ati
ng
a b
usi
ne

ss i
de
a u
sin
g t
he
Thr
ee-
Filt
er
Te
st.
He
re
is t
he
ide
a: [
de
scr
ipti
on]
. H
elp
m
e a
ss
es
s w
het
het
her
it
pa
ss
es
ea
ch
of t
he
thr
ee
filt
ers
: A
gre
em
ent
(d
o c
ust
om
ers
in

de
pe
nd
ent
ly
na
me
thi
s p
ain
?),
Sc
ale
(d
oe
s t
he
ma
rke
t s
up
por
t a
rea
l b
usi
ne
ss
?),
an
d
Co
mp
ete
nc
e (
do
es
my
te
am
ha
ve
unf
air
ad
va
nta
ge
s?)
. Id
ent
ify

the
w
ea
ke
st f
ilte
r a
nd
su
gg
est
ho
w
to
ad
dre
ss
it."

Task 3: Jobs to Be Done Interview (60 min)

- Id
ent
ify
on
e p
ers
on
wh
o h
as
rec
ent
ly
pur
ch
as
ed
or
sto
pp
ed
usi
ng
a p
rod
uct
in
yo
ur t
arg
et
do

ma
in.

- C
on
du
ct
a
30-
mi
nut
e h
ire
d
or f
ire
d i
nte
rvi
ew
us
ing
th
e s
cri
pt
in
Se
cti
on
2.3
.

- A
fter
th
e i
nte
rvi
ew,
wr
ite
a 2
00-
wo
rd
sy
nth
esi
s t
hat
an
sw
ers

:
(a)
What was the triggering event?
(b)
What was the functional job?
(c)
What was the emotional job?
(d)
What was the social job?
• Use this starter

prompt (after meta-prompting): "I conducted a customer interview and here are my notes: [past notes]. Help me identify the functional job, e motivational job,

and social job the customer was trying to accomplish. What does this suggest about how we should frame the evaluation proposition?

Task 4: Value Proposition Canvas (60 min)

- Download the V

PC
te
mp
lat
e fr
om
St
rat
eg
yz
er
or r
ecr
eat
e
it.

- C
om
ple
te t
he
Cu
sto
me
r P
rofi
le
sid
e:
list
5+
cu
sto
me
r jo
bs,

5+
cu
sto
me
r p
ain
s (r
an
ke
d
by
int
en
sity
, a
nd

3+
cu
sto
me
r g
ain
s.

- C
om
ple
te t
he
Val
ue
Ma
p s
ide
:
list
yo
ur
pro
po
se
d p
rod
uct
/se
rvi
ce,
id
ent
ify
whi
ch
pai
ns
it r
eli
ev
es
an
d h
ow,
an
d i
de
ntif
y w
hic
h g
ain
s

it c
rea
tes

.

- A
chi
ev
e
fit:
cir
cle
th
e p
ain
s a
nd
gai
ns
on
the
C
ust
om
er
Pr
ofil
e t
hat
yo
ur
Val
ue
Ma
p d
ire
ctly
ad
dre
ss
es.
A
str
on
g V
PC
ad
dre
ss
es
at l
ea
st
3

of the
top-5
pairs.

- Use this
starter
prompt (after
meta-
prompting):
"

Here is my
completed
Value
Proposition
Canvas:
[describe
both sides].
Evaluate
the

e s
tre
ngt
h
of t
he
fit
bet
we
en
my
va
lue
m
ap
an
d c
ust
om
er
pro
file
. Id
ent
ify
the
thr
ee
big
ge
st
ga
ps.
Su
gg
est
ho
w
I c
oul
d r
ed
esi
gn
the
va
lue
pr
op
osi
tio
n
to

ac
hie
ve
str
on
ger

fit."

Task 5: Moat Analysis (45 min)

- R
evi
ew
th
e f
our

Al-
nat
ive
m
oat
s fr
om
Se
cti
on
2.4
: W
ork
flo
w
De
pth
, P
rop
riet
ary
Fe
ed
ba
ck
Lo
op
s,
Ag
ent
Gr
ap
h P
osi
tio

n,
an
d P
ro
mp
t
IP.

• F
or
yo
ur
pro
po
se
d s
olu
tio
n,
wri
te
on
e p
ara
gra
ph
per
m
oat
:
(a)
Do
es
thi
s
mo
at
ap
ply
to
yo
ur
sol
uti
on
?
(b)
Ho
w
wo
uld
yo
u d
ev

elo
p
it?
(c)
Wh
at
is t
he
tim
eli
ne
to
bui
ld
me
ani
ngf
ul
de
pth
?

• U
se
thi
s s
tart
er
pro
mp
t (a
fter
m
eta
-pr
om
pti
ng)
: "I'
m
an
aly
zin
g t
he
co
mp
etit
ive
m
oat
s
for
thi

s
Al-
nat
ive
pr
od
uct
: [d
es
cri
pti
on]
. E
val
uat
e
whi
ch
of t
he
fou
r
Al-
nat
ive
m
oat
s (
wo
rkfl
ow
de
pth
, pr
opr
iet
ary
fe
ed
ba
ck l
oo
ps,
ag
ent
gr
ap
h p
osi
tio
n,
pro
mp

t
IP
an
d e
val
s)
ap
ply
to
thi
s p
rod
uct
. F
or t
he
tw
o s
tro
ng
est
, d
es
cri
be
sp
eci
fic
ally
ho
w t
he
co
mp
an
y s
ho
uld
in
ve
st
to
bui
ld t
ho
se
mo
ats
ov
er t
he
ne
xt

AI Studio Build — Weekly Application

The Pain Translator

Every business problem starts as an abstraction. "Supply chain inefficiency." "Customer churn." "Employee engagement." These abstractions are analytically inert — they tell you nothing about who is suffering, when, or how much. The Pain Translator is an AI Studio application that converts abstract problem statements into structured, actionable pain narratives.

What you are building: In aistudio.google.com, build a structured-output prompt that takes an abstract problem statement (e.g., "inventory challenges for small retailers") and returns a JSON object containing:

- p
ers
on
a:
a n
am
ed,
sp
eci
fic
per
so
n
wit
h
a r
ole
, c
om
pa
ny
typ
e,
an
d l
oc
ati
on
—
not
a
de

mo
gra
phi
c c
ate
gor
y

• p
ain
_m
om
ent
: th
e s
pe
cifi
c
mo
me
nt
of
ac
ute
pa
in
— i
ncl
ud
e t
he
tim
e,
pla
ce,
w
hat
ha
pp
en
ed,
an
d
wh
o
wa
s a
ffe
cte
d

• q
ua
ntif

ied
_c
ost
: a
co
ncre
ete
m
ea
sur
e
of t
he
co
st
—
tim
e l
ost
in
ho
urs
, d
oll
ars
lo
st,
or
de
als
mi
ss
ed
• e
mo
tio
nal
_c
on
se
qu
en
ce:
ho
w t
he
per
so
na
felt
in
tha
t m

om
ent
an
d
in t
he
ho
urs
aft
er

—
sp
eci
fic
em
oti
on
al l
an
gu
ag
e,
not
ge
ner
ic
de
scr
ipt
ors

• ro
ot_
ca
us
e: t
he
str
uct
ura
l re
as
on
wh
y t
his
pr
obl
em
oc
cur
red

—

wh
at
org
ani
zat
ion
al,
tec
hni
cal
,
or
ma
rke
t fa
ilur
e c
rea
ted
th
e c
on
diti
on
s

Step-by-step:

- O
pe
n a
ist
udi
o.g
oo
gle
.co
m
an
d n
avi
gat
e
to
Bui
ld.
- S
ele
ct
a
Ge
mi
ni

model and enable JS ON mode (or "Structured output put")

— the toggle is in the model settings panel).

- Define a response schema that enforces the

e-fi
eld
str
uct
ure
ab
ov
e.
In
Al
Stu
dio
, th
is
is
do
ne
by
clic
kin
g "
Edi
t S
ch
em
a"
an
d d
efi
nin
g t
he
JS
ON
sc
he
ma
:

```
    "pain_moment": { "type": "string" },  
    "quantified_cost": { "type": "string" },  
    "emotional_consequence": { "type": "string" },  
    "root_cause": { "type": "string" }  
  },  
  "required": ["persona", "pain_moment", "quantified_cost", "emotional_consequence",  
    "root_cause"]  
}
```

te
a s
yst
em

in
str
uct
ion
th
at t
ells
th
e
mo
del
:

(a)
its
rol
e
as
a b
usi
ne
ss
pai
n a
nal
yst

,
(b)
the
le
vel
of
sp
eci
fikit
y r
eq
uir
ed
(na
me
d p
ers
on,
sp
eci
fic
mo
me
nt,
qu
ant
ifie
d c

ost
, a
nd
(c)
the
qu
alit
y b
ar (
out
put
s
mu
st
be
vis
cer
al
an
d c
on
cre
te,
not
ge
ner
ic
or
ab
str
act
) . A
ppl
y
me
ta-
pro
mp
tin
g
to
de
vel
op
a t
hor
ou
gh
2—
3 p
ara
gra
ph

sys
te
m i
nst
ruc
tio
n.

• T
est
yo
ur
Pai
n T
ran
sla
tor
wit
h
at l
ea
st t
hre
e d
iffe
ren
t a
bst
rac
t pr
obl
em
st
ate
me
nts
fro
m t
hre
e d
iffe
ren
t in
du
stri
es:

-
A r
eta
il o
per
ati
on
s p

rob
le
m

-

A h
eal
thc
are
op
era
tio
ns
pro
ble
m

-

A fi
na
nci
al
ser
vic
es
or
pro
fes
sio
nal
se
rvi
ce
s p
rob
le
m

- F
or
ea
ch
tes
t, e
val
uat
e t
he
out
put
:
Is t
he
per
so
na

specific and believable? Is the pain moment vivid? Is the cost quantified? Refine your system instruction based on what you observe.

- Use the Share feature

nct
ion
to
ge
ner
ate
a
pu
blic
AI
Stu
dio
sh
are
lin
k.

Deliverable: Submit (a) your AI Studio share link, (b) the full JSON schema you defined, and (c) three sample Pain Translator outputs for three different problem statements. Your system instruction should be 2–3 substantial paragraphs.

What you are learning: Structured output, JSON mode, response schema definition. This is the foundational capability for any AI application that needs to produce machine-readable output rather than freeform text — essential for every subsequent AI Studio Build assignment and for any production AI pipeline you will build.

Discussion Prompts

Discussion 1: The Pain-First Inversion in Your Industry

Select a product or service in your industry that you believe was built "features first" — designed around a technical capability before a validated customer pain. Using the frameworks from this chapter, reconstruct what a "pain-first" version of that product's development might have looked like.

In your post:

- Id
ent
ify
the
sp
eci
fic
fea
tur
e
or
ca
pa
bilit

y t
he
co
mp
an
y l
ed
wit
h,
an
d
wh
at
cu
sto
me
r p
ain
yo
u b
eli
ev
e
it s
ho
uld
ha
ve
led
wit
h

- D
es
cri
be
wh
at t
he
"S
ara
h v
ers
ion
"

of t
hat
pa
in
wo
uld
lo
ok
like

—
a n
am
ed
per
so
na,
a
sp
eci
fic
ba
d
mo
me
nt,
a q
ua
ntif
ied
co
st

- A
ppl
y t
he
Thr
ee-
Filt
er
Te
st:
do
es
the
pa
in
yo
u i
de
ntif
ied
pa
ss
all
thr
ee
filt
ers
?

- A

rgu
e
wh
eth
er t
he
co
mp
an
y's
fea
tur
es-
firs
t a
ppr
oa
ch
wa
s
a s
trat
egi
c e
rro
r
or
a p
rag
ma
tic
pro
du
ct
de
cisi
on
giv
en
ma
rke
t c
on
diti
on
s
at t
he
tim
e

Support your argument with at least two scholarly sources on value proposition design, customer development, or innovation strategy. Include at least one current source (no more than two years

old) documenting outcomes for the company you selected.

Discussion 2: The Moat That Isn't a Moat

Many AI companies currently claim competitive advantages that are, under careful analysis, not durable moats — they are first-mover advantages that will erode as model capabilities improve and competitive intensity increases.

In your post:

- Select a real AI company (public or well-documented private) and identify the competitive advantage the company

ntl
y cl
ai
m

- A
ppl
y t
he
Mo
at
Du
rab
ility
M
atri
x fr
om
Se
cti
on
2.4
: is
thi
s
a d
ura
ble
m
oat
or
an
ero
din
g o
ne
?

- Id
ent
ify
whi
ch
of t
he
fou
r
Al-
nat
ive
m
oat
s (
wo

rkfl
ow
de
pth
, fe
ed
ba
ck l
oo
ps,
ag
ent
gr
ap
h p
osi
tio
n,
pro
mp
t IP
) th
e c
om
pa
ny
is
mo
st li
kel
y
to
de
vel
op
ov
er t
he
ne
xt t
hre
e y
ear
s
—
or
arg
ue
tha
t th
ey
will
no

t d
ev
elo
p
a d
ura
ble
m
oat
at
all

• M
ak
e
a p
red
icti
on:
wil
l th
is c
om
pa
ny
still

be
am
on
g t
he
top
thr
ee
pla
yer
s
in
its
cat
eg
ory
in
thr
ee
ye
ars
?
De
fen
d y
our
pr

editi
on
with
the
vid
en
ce.

Your post should demonstrate familiarity with at least one of the following scholarly frameworks: Porter's Five Forces (Porter, 1980), competitive dynamics (D'Aveni, 1994), or resource-based view of the firm (Barney, 1991).

Readings

Required:

- O
ste
rw
ald
er,
A.,
Pig
ne
ur,
Y.,
Be
rna
rda
, G
,
&
Sm
ith,
A.
(20
14)
. V
alu
e p
rop
osi
tio
n d
esi
gn:
H
ow
to

create products and services customers want. Wiley. (Chapter 1–2: The Value Proposition Canvas)
• Christensen, C. M., Hall, T., Dillon, K., & Duncan

an,
D.
S. (
20
16)
. K
no
w y
our
cu
sto
me
rs'
"jo
bs
to
be
do
ne.
" H
arv
ard
Bu
sin
es
s R
evi
ew,
94
(9),
54
—6
2.
Av
ail
abl
e
at
hbr
.or
g
• Fi
tzp
atri
ck,
R.
(20
13)
. T
he
Mo
m
Te

st:
Ho
w
to t
alk
to
cu
sto
me
rs
an
d l
ear
n
if y
our
bu
sin
es
s
is
a g
oo
d i
de
a
wh
en
ev
ery
on
e
is l
yin
g
to
yo
u.
Ro
bfit
z L
td.
(C
ha
pte
rs
1–
3)
• P
ort
er,
M.
E. (

19
80)
. C
om
pet
itiv
e s
trat
eg
y:
Te
ch
niq
ue
s
for
an
aly
zin
g i
nd
ust
rie
s a
nd
co
mp
etit
ors
. Fr
ee
Pr
es
s. (
Ch
apt
er
1:
Th
e S
tru
ctu
ral
An
aly
sis
of I
nd
ust
rie
s)

Recommended:

- Christensen, C. M., & Raynor, M. E. (2003). The innovator's solution: Creating and sustaining successful growth. Harvard Business Review Press.
- Moore

sta

,

B. (

20

23)

. D

em

an

d-s

ide

sa

les

10

1:

Sto

p s

elli

ng

an

d h

elp

yo

ur

cu

sto

me

rs

ma

ke

pro

gre

ss.

Lio

ncr

est

Pu

blis

hin

g.

• B

arn

ey,

J. (

19

91)

. Fi

rm

res

our

ce

s a

nd

su

sta
ine
d c
om
pet
itiv
e a
dv
ant
ag
e.
Jo
urn
al
of
Ma
na
ge
me
nt,
17(
1),
99
–1
20.
• D
'Av
eni
, R
.A.
(19
94)
. H
yp
erc
om
pet
itio
n: Ma
na
gin
g t
he
dy
na
mi
cs
of
str
ate
gic
m

an
eu
ver
ing
. Fr
ee
Pr
es
s.

• Bl
an
k,
S.,
&
Do
rf,
B. (2012).
The
sta
rtu
p o
wn
er's
ma
nu
al:
Th
e s
tep
-
by-
ste
p g
uid
e
for
bui
ldi
ng
a g
rea
t c
om
pa
ny.
K&
S
Ra
nc

h P
res
s.

Glossary

Term	Definition
Value Proposition	The specific combination of benefits a product or service delivers to a defined customer segment.
Pain-First Framework	A product development and pitch methodology that prioritizes identifying and articulating customer pain points before benefits.
The Sarah Exercise	A technique for transforming abstract, category-level pain descriptions into specific, vivid persona narratives.
Pain Quantification	The practice of measuring a customer pain across four dimensions: time wasted, dollars lost, deals missed, and customer frustration.
Three-Filter Test	A structured framework for evaluating business opportunities across three criteria: Agreement (do customers have the problem?), Scale (affects enough people?), and Advantage (can we solve it better?).
Agreement Filter	The first filter in the Three-Filter Test: a problem passes when it affects at least 10,000 addressable customers.
Scale Filter	The second filter in the Three-Filter Test: a problem passes when it affects at least 10,000 addressable customers.
Competence Filter	The third filter in the Three-Filter Test: a problem passes when the team has unfair, asymmetric advantages to solve it.
Jobs to Be Done (JTBD)	Clayton Christensen's framework holding that customers do not buy products — they hire solutions to take specific jobs in a restrictive environment to achieve the most visible dimension of ITRD.
Functional Job	One of the three job types in JTBD: the job that achieves the most visible dimension of ITRD.
Emotional Job	How the customer wants to feel after hiring a product; the motivating dimension of JTBD that determines willingness to pay.
Social Job	How the customer wants to be perceived by others after using a product; particularly important in B2B.
Hired and Fired Interview	A structured qualitative technique that reconstructs purchase stories (hired) and abandonment stories (fired) to understand competitive advantages.
Competitive Moat	A durable structural advantage that allows a company to generate above-market returns for an extended period.
Workflow Depth	An AI-native moat describing how deeply an AI product is embedded in the customer's daily operational fabric.
Proprietary Feedback Loop	A reinforcing loop in which usage generates training signals that improve the AI system, creating a data advantage.
Agent Graph Position	An AI-native moat derived from occupying the orchestration or coordination layer in a multi-agent AI system.
Prompt IP	A systematic approach to curating and accumulating prompt libraries, evaluation frameworks, and fine-tuning datasets.
Value Proposition Canvas (VPC)	A strategic design tool by Osterwalder and Pigneur that maps customer jobs, pains, and gains against a company's offerings.
Capability Commoditization	The process by which previously scarce capabilities become widely available, eroding moats built on model access rather than structural position.